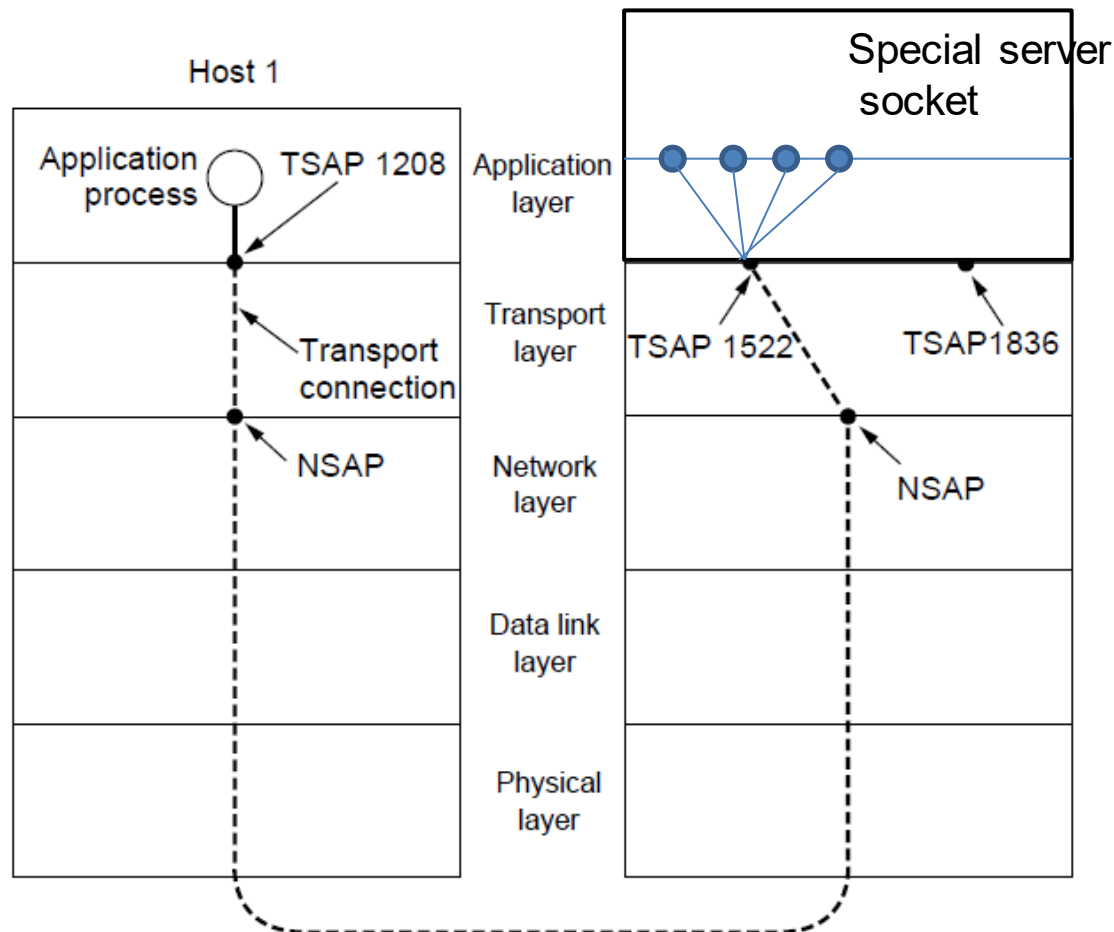

Transport Layer Contd

Internet Technologies
COMP90007

Sockets & Addressing

- Socket library provides a type of multiplexing tool on top of TSAPs to allow servers to service multiple clients



Port Allocations

- Port numbers can range from 0-65535
- Port numbers are regulated by IANA (<http://www.iana.org/assignments/port-numbers>)
- Ports are classified into 3 segments:
 - Well Known Ports (0-1023)
 - Registered Ports (1024-49151)
 - Dynamic Ports (49152-65535)

Port	Protocol	Use
20, 21	FTP	File transfer
22	SSH	Remote login, replacement for Telnet
25	SMTP	Email
80	HTTP	World Wide Web
110	POP-3	Remote email access
143	IMAP	Remote email access
443	HTTPS	Secure Web (HTTP over SSL/TLS)
543	RTSP	Media player control
631	IPP	Printer sharing

Transport Layer Protocols: UDP


- The **most basic service** is actually connectionless:
 - Called: User Datagram Protocol (UDP)
 - Does not add much to the Network Layer functionality
 - TCP gives the real-deal for this layer, *reliability*...
 - For UDP: Just remove connection primitives to use it in a program
- **UDP good for?:**
 - It is used for apps like video streaming/gaming regularly
- **The reliability issue is left to?:**
 - the application layer... retransmission decisions as well as congestion control

UDP Contd

- Provides a protocol whereby **applications can transmit encapsulated IP datagrams without a connection establishment**
- UDP transmits in segments consisting of an **8-byte header followed by the payload**
- UDP **headers contain source and destination ports**
- Payload is handed to the process which is attached to the particular port at destination

UDP Contd.

- Main **advantage** of using UDP over raw IP is:
 - the ability to specify ports for source and destination pairs, i.e., **addressing for processes**
- Both source and destination ports are required - destination allows for incoming segments, source allows reply for outgoing segments



Structure of UDP header: It has ports (TSAPs), length and checksum

Strengths and Weaknesses of UDP

- **Strengths:** provides an IP interface with multiplexing/de-multiplexing capabilities and related transmission efficiencies
- **Weaknesses:** UDP does not include support for: flow control, error control/retransmission of bad segments
- **Conclusion:** where applications require a precise level of control over packet flow/error/timing, UDP is a good choice as application layer can make choices
- **Domain Name System over the Internet is a famous user of UDP**


Another Application of UDP: Remote Procedure Call (RPC)

- Sending a message and getting a reply back is analogous to making **a function call** in programming languages
- Birrell and Nelson modified this to allow programs to call procedures on remote hosts using UDP
 - **Remote Procedure Call (RPC)**

Remote Procedure Call (RPC)

- To call a remote procedure, the client is bound to a small library (the **client stub**) that represents the server procedure in the client's address space.
- Similarly the server is bound with a procedure called the **server stub**.
- These **stubs hide the fact that the procedure itself is not local.**

RPC Illustrated



Transmission Control Protocol:

TCP Details

- Provides a protocol by which applications can transmit datagrams within a **connection-oriented** framework, thus increasing reliability
- TCP transport entity manages TCP streams and interfaces to the IP layer - can exist in numerous locations (kernel, library, user process)
- **TCP entity** accepts user data streams, and **segments them into pieces < 64KB** (often at a size in order so that the IP and TCP headers can fit into a single Ethernet frame), and sends each piece as a separate datagram
- Recipient TCP entities reconstruct the original byte streams from the encapsulation


The TCP Service Model

- Sender and receiver both create **sockets**, consisting of the IP address of the host and a port number as we saw earlier
- For TCP Service to be activated, **connections must be explicitly established between a socket at a sending host** (src-host, src-port) and a socket at a receiving host (dest-host, dest-port)
- Special one-way server sockets may be used for multiple connections simultaneously

Example

Host 128.42.11.3

Host 62.118.44.12



Features of TCP Connections

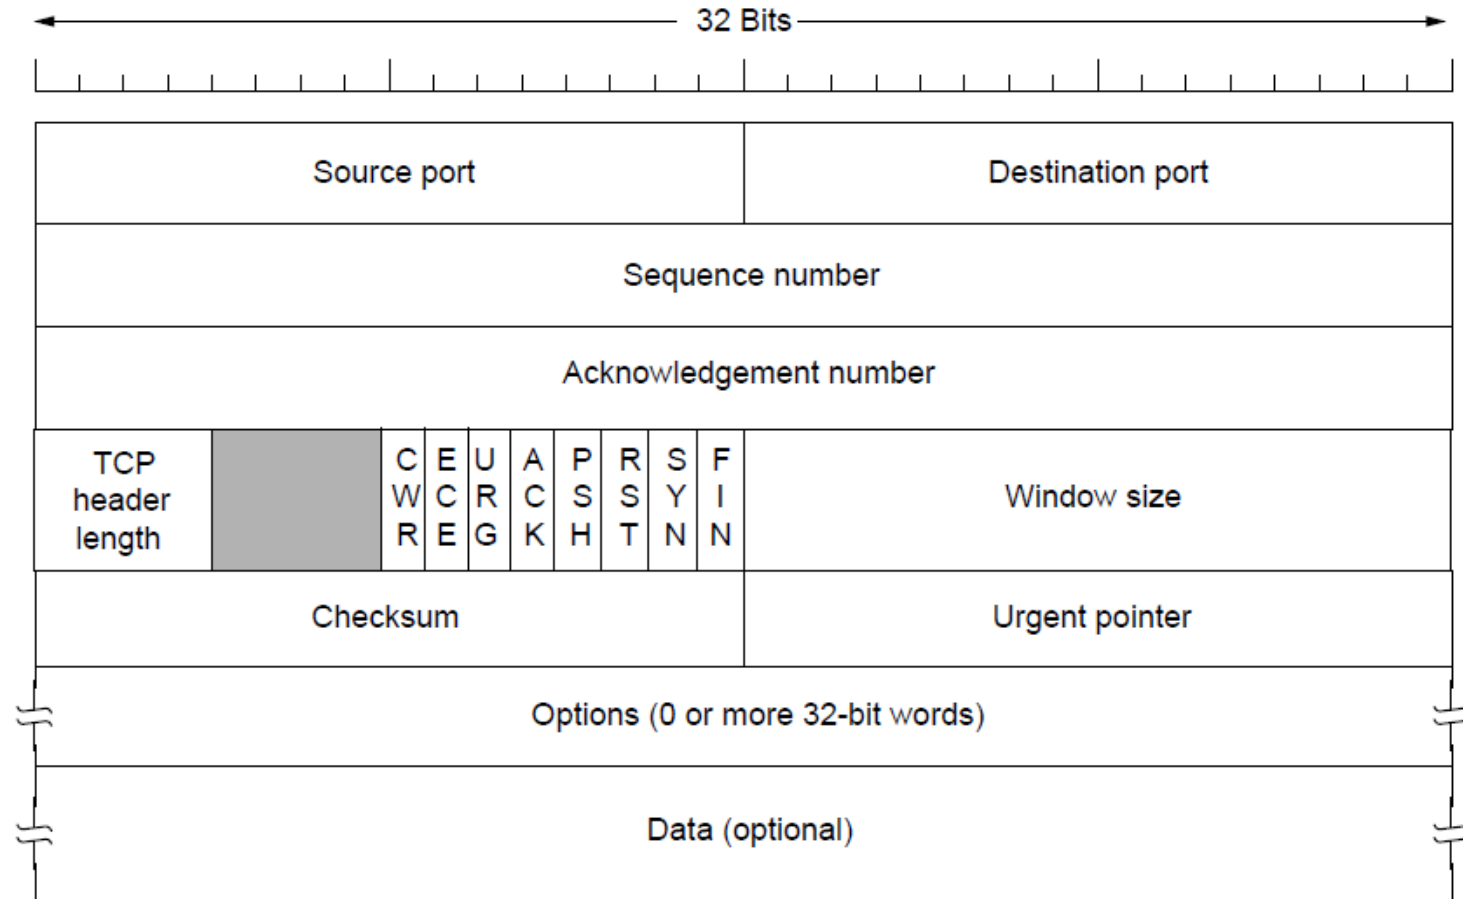
- TCP connections are:
- **Full duplex** - data in both directions simultaneously
- **Point to point** - exact pairs of senders and receivers
- **Byte streams**, not message streams - message boundaries are not preserved
- **Buffer options** - TCP entity can choose to buffer prior to sending or not depending on the context
 - **TCP_NODELAY in Java**
 - **Socket.setTcpNoDelay(boolean)**

TCP Contd

- Data sent between TCP entities in segments - segment has a **20 byte header plus zero or more data bytes**
- TCP entities decide how large segments should be mainly with 2 constraints:
 - 65,515 byte IP payload
 - Ethernet unit size - generally 1500 bytes
- **Sliding window** - **sender transmits and starts a timer**
 - Receiver sends back an acknowledgement which is the next sequence number expected - if sender's timer expires before acknowledgement, then the sender **transmits the original segment again**

The TCP Segment Header

- TCP header includes addressing (ports), sliding window (seq. / ack. number), flow control (window), error control (checksum) and more



The TCP Segment Header

- **Source port** and **Destination port** fields identify the local end points of the connection
- **Sequence number** and **Acknowledgement number** fields perform their usual functions
- **TCP header length** tells how many 32-bit words are contained in the TCP header
- **Window size** field tells how many bytes may be sent starting at the byte acknowledged
- **Checksum** is also provided for extra reliability. It checksums the header, the data
- **Options** field provides a way to add extra facilities not covered by the regular header
- **URG** is set to 1 if the *Urgent pointer* is in use. The Urgent pointer is used to indicate a byte offset from the current sequence number at which urgent data are to be found

The TCP Segment Header

- **CWR** and **ECE** are used to signal congestion when *ECN* (Explicit Congestion Notification) is used
- **ECE** is set to signal an ECN-Echo to a TCP sender to tell it to slow down when the TCP receiver gets a congestion indication from the network
- **CWR** is set to signal Congestion Window Reduced from the TCP sender to the TCP receiver so that it knows the sender has slowed down and can stop sending the ECN-Echo
- The **ACK** bit is set to 1 to indicate that the Acknowledgement number is valid. This is the case for nearly all packets. 0 means ignore ACK number field
- **PSH** bit indicates PUSHed data. The receiver is hereby kindly requested to deliver the data to the application upon arrival and not buffer it until a full buffer has been received

The TCP Segment Header

- The **RST** bit is used to abruptly reset a connection that has become confused due to a host crash or some other reason. It is also used to reject an invalid segment or refuse an attempt to open a connection
- The **SYN** bit is used to establish connections. The connection request has $\text{SYN} = 1$ and $\text{ACK} = 0$. The connection reply does bear an acknowledgement, so it has $\text{SYN} = 1$ and $\text{ACK} = 1$.
- In essence, the SYN bit is used to **denote both CONNECTION REQUEST and CONNECTION ACCEPTED**, with the ACK bit used to distinguish between those two possibilities.
- The **FIN** bit is used to release a connection. It specifies that the sender has no more data to transmit. However, after closing a connection, the closing process may continue to receive data.