

# COMP90038 Algorithms and Complexity

## Balanced Trees

Douglas Pires

Lecture 15

Semester 1, 2021

# Approaches to Balanced Binary Search Trees

To optimise the performance of BST search, it is important to keep trees (reasonably) balanced.

Instance simplification approaches: Self-balancing trees

- AVL trees
- Red-black trees
- Splay trees

Representational changes:

- 2-3 trees
- 2-3-4 trees
- B-trees

# AVL Trees

Named after Adelson-Velsky and Landis.





Recall that we defined the height of the empty tree as -1.

For a binary (sub-) tree, let the **balance factor** be the difference between the height of its left sub-tree and that of its right sub-tree.

An **AVL tree** is a BST in which the balance factor is -1, 0, or 1, for every sub-tree.

# AVL Trees: Examples and Counter-Examples

Which of these are AVL trees?




# Building an AVL Tree

As with standard BSTs, insertion of a new node always takes place at the fringe of the tree.

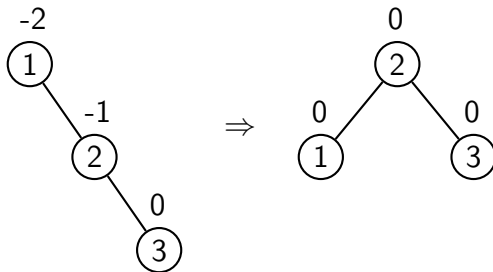
If insertion of the new node makes the AVL tree unbalanced (some nodes get balance factors of 2 or -2), transform the tree to regain its balance.

Regaining balance can be achieved with one or two simple, local transformations, so-called **rotations**.

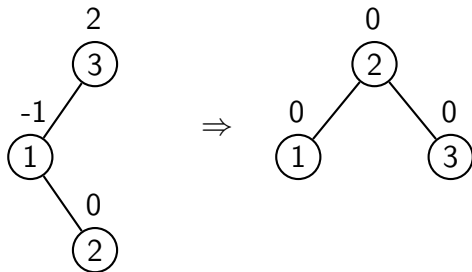
# AVL Trees: R-Rotation



# AVL Trees: L-Rotation

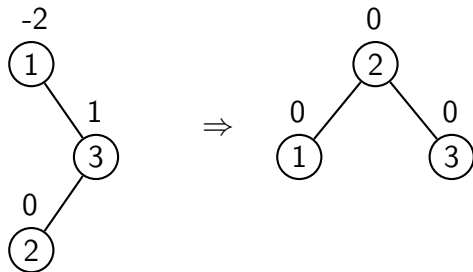


# AVL Trees: LR-Rotation



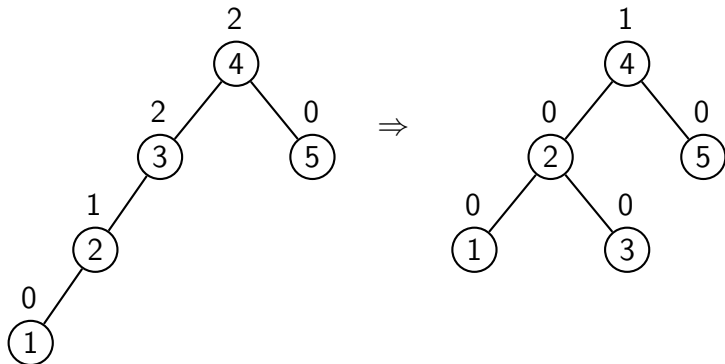


# AVL Trees: RL-Rotation



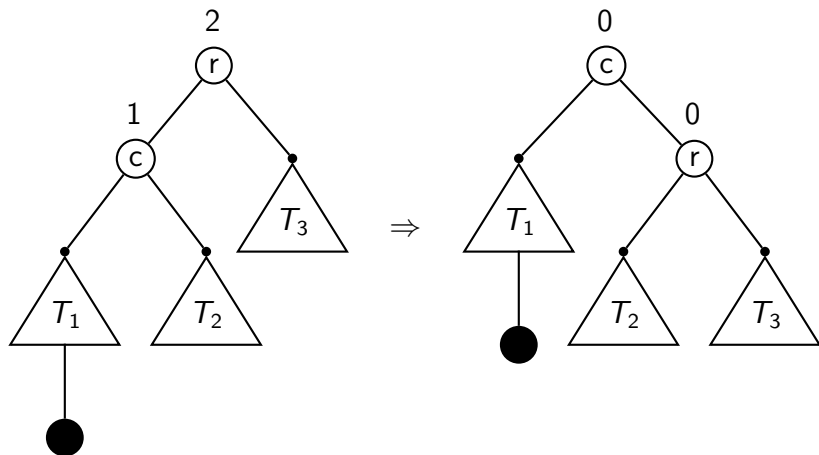
# AVL Trees: Where to Perform the Rotation

Along an unbalanced path, we may have several nodes with balance factor 2 (or -2):



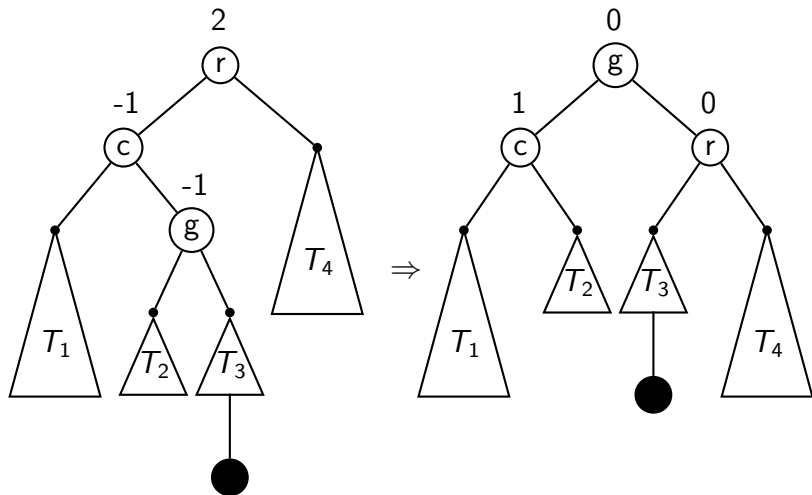
It is always the **lowest** unbalanced subtree that is re-balanced.

# AVL Trees: The Single Rotation, Generally



This shows an **R-rotation**; an **L-rotation** is similar.

# AVL Trees: The Double Rotation, Generally



This shows an **LR-rotation**; an **RL-rotation** is similar.

# Properties of AVL Trees

The notion of “balance” that is implied by the AVL condition is sufficient to guarantee that the depth of an AVL tree with  $n$  nodes is  $\Theta(\log n)$ .

For random data, the depth is very close to  $\log_2 n$ , the optimum.


In the worst case, search will need at most 45% more comparisons than with a perfectly balanced BST.

**Deletion** is harder to implement than insertion, but also  $\Theta(\log n)$ .

# Other Kinds of Balanced Trees

A **red-black tree** is a BSTs with a slightly different concept of “balanced”. Its nodes are coloured red or black, so that

- 1 No red node has a red child.
- 2 Every path from the root to the fringe has the same number of black nodes.



A worst-case red-black tree (the longest path is twice as long as the shortest path).

A **splay tree** is a BST which is not only self-adjusting, but also **adaptive**. Frequently accessed items are brought closer to the root, so their access becomes cheaper.

## 2-3 Trees

2-3 trees and 2-3-4 trees are search trees that allow more than one item to be stored in a tree node.

As with BSTs, a node that holds a single item has (at most) two children.

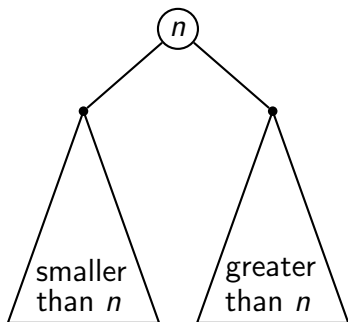
A node that holds two items (a so-called **3-node**) has (at most) three children.

And for 2-3-4 trees, a node that holds three items (a **4-node**) has (at most) four.

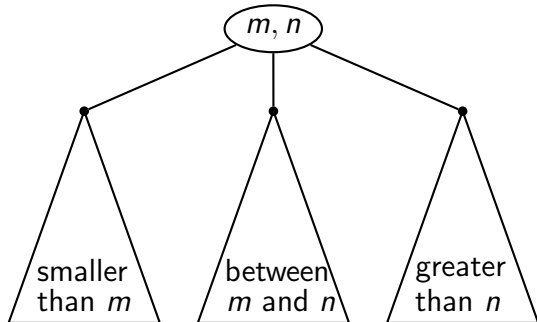
This allows for a simple way of keeping search trees balanced.

# 2-Nodes and 3-Nodes

2-node



3-node





# Insertion in a 2–3 Tree

To insert a key  $k$ , pretend that we are searching for  $k$ .



This will take us to a leaf node in the tree, where  $k$  should now be inserted; if the node we find there is a 2-node,  $k$  can be inserted without further ado.

Otherwise we had a 3-node, and the two inhabitants, together with  $k$ , momentarily form a node with three elements; in sorted order, call them  $k_1$ ,  $k_2$ , and  $k_3$ .

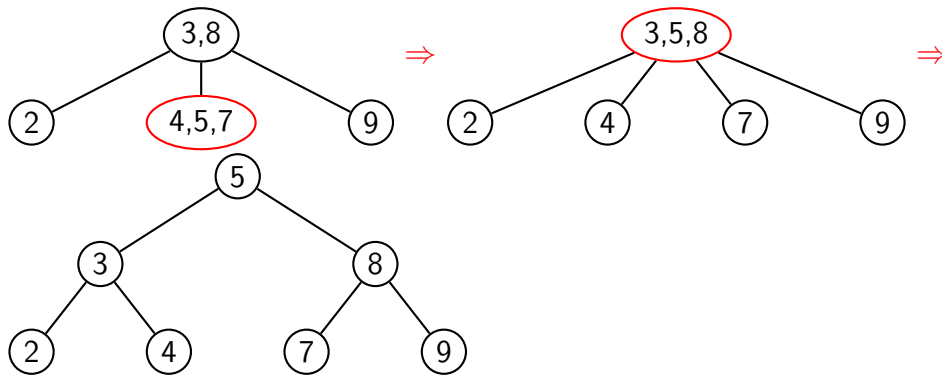
We now **split** the node, so that  $k_1$  and  $k_3$  form their own individual 2-nodes. The middle key,  $k_2$  is **promoted** to the parent node.

The promotion may cause the parent node to overflow, in which case **it** gets split the same way. The only time the tree's height changes is when the root overflows.

# Example: Build a 2-3 Tree from 9, 5, 8, 3, 2, 4, 7



# Example: Build a 2-3 Tree from 9, 5, 8, 3, 2, 4, 7



# Exercise: 2–3 Tree Construction

Build the 2–3 tree that results from inserting these keys, in the given order, into an initially empty tree:

C, O, M, P, U, T, I, N, G



## 2-3 Tree Analysis

Worst case search time results when all nodes are 2-nodes.

The relation between the number  $n$  of nodes and the height  $h$  is:

$$n = 1 + 2 + 4 + \dots + 2^h = 2^{h+1} - 1$$

That is,  $\log_2(n + 1) = h + 1$ .

In the best case, all nodes are 3-nodes:

$$n = 2 + 2 \times 3 + 2 \times 3^2 + \dots + 2 \times 3^h = 3^{h+1} - 1$$

That is,  $\log_3(n + 1) = h + 1$ .

Hence we have  $\log_3(n + 1) - 1 \leq h \leq \log_2(n + 1) - 1$ .

Useful formula:  $\sum_{i=0}^n a^i = \frac{a^{n+1}-1}{a-1}$ , for  $a \neq 1$ .

# Coming Soon to a Theatre Near You

How to buy time, by spending a bit of space.