

Lecture 1: Introduction and Overview

COMP90049

Introduction to Machine Learning

Semester 2, 2021

Lida Rashidi, CIS

Copyright @ University of Melbourne 2021. All rights reserved. No part of the publication may be reproduced in any form by print, photoprint, microfilm or any other means without written permission from the author.

Acknowledgement: Lea Frermann



This lecture

- Introduction and Warm-up
- Housekeeping COMP90049
- Machine Learning

Intros & Warm-up

About Lida

- Lecturer in CIS since 2019
- Research in graph mining and information retrieval
- PhD from the University of Melbourne
- 4 years of research in academia

About Qiuhong

- Lecturer in CIS since 2020
- Research in machine learning and computer vision
- PhD from the University of Western Australia
- 1.5 years of research in Max Planck Institute for Informatics, Germany



About Lida

- Lecturer in CIS since 2019
- Research in graph mining and information retrieval
- PhD from the University of Melbourne
- 4 years of research in academia

About Qiuhong

- Lecturer in CIS since 2020
- Research in machine learning and computer vision
- PhD from the University of Western Australia
- 1.5 years of research in Max Planck Institute for Informatics, Germany

About you

Please go to: pollev.com/comp90049



What is Learning?

What is Machine Learning?

Some proposed definitions...

“The computer automatically learns something”

“Statistics, plus marketing”

“ ... how to construct computer programs that automatically improve with experience A computer program is said to learn from experience ... if its performance ... improves with experience... ”

Mitchell [1997, pp. xv-17]

“We are drowning in information, but we are starved for knowledge”

John Naisbitt, Megatrends

Our definition of Machine Learning

automatic extraction of **valid, novel, useful** and **comprehensible knowledge** (rules, regularities, patterns, constraints, models, ...) from arbitrary sets of data

Learning what?

- **Task** to accomplish a goal, e.g.,
 - Assign continuous values to inputs (essay \rightarrow grade)
 - Group inputs into known classes (email \rightarrow {spam, no-spam})
 - Understand regularities in the data

Learning from what?

- **Data**
- Where do the data come from? Is it reliable? Representative?

How do we learn?

- define a **model** that explains how to get from input to output
- derive a **learning algorithm** to find the best model parameters

How do we know learning is happening?

- The algorithm improves at its task with exposure to more data
- We need to be able to **evaluate** performance objectively



About COMP90049

Coordinator & Lecturer1	Lida Rashidi	rashidil@unimelb.edu.au
Lecturer 2	QiuHong Ke	quihong.ke@unimelb.edu.au
Tutors	Tahrima Hashem	tahrima@unimelb.edu.au
	Pei-Yun Sun	pssun@unimelb.edu.au
	Ella Alipourchavary	ella.alipourchavary@unimelb.edu.au
	Kazi Adnan	kazi.adnan@unimelb.edu.au
	Hasti Samadi	hasti.samadi@unimelb.edu.au
	Zenan Zhai	zenan.zhai@unimelb.edu.au

- The lectures will be delivered **fully online**
- I'll aim for as much interaction as possible (and desired)
- All live lectures will be recorded. All recordings and other materials will be made available online through Canvas
- **Live lectures** via Zoom for the first couple of weeks
- Afterwards possibly **pre-recorded with live Q&A sessions**
- **Live** and **in-person** workshops throughout the semester **dual delivery mode**

Lectures

Lecture 1	Wed 16:15-17:15 Online; Zoom
Lecture 2	Thu 14:15-15:15 Online; Zoom

Lecture content

- Theory
- Derivation of ML algorithms from scratch
- Motivation and context
- Some coding demos in Python



Workshops

- **start from week 2**
- 1 hour per week
- ~ 14 slots, please sign up and stick to one
- Online workshops are live via zoom and In-person workshops will be on campus
- At the moment, due to restrictions the on campus workshops are converted to online workshops
- Return to campus will be announced on LMS

Workshop Content

- Practical exercises
- Working through numerical examples
- Revising theoretical concepts from the lectures



Coding drop-in sessions

Session 1	Tuesday 12:00–13:00 (link via Canvas Zoom)
Session 2	Friday 15:00–16:00 (link via Canvas Zoom)

- **start from week 2** and run until week 5
- you can ask questions around Python / the weekly code snippets
- **Not** an assignment consultation

Materials and announcements

- All materials will be made available through LMS (Canvas)
- Important news will be shared via Canvas Announcements (expect about 1 per week)

General inquiries: Piazza forum on LMS

- We encourage all students to join in discussions answering other students questions is one of the best ways to improve your own understanding
- Please do not post sections of your code or reports publicly on Piazza! If you must include these, private-message the instructors

Personal/private concerns: Email your tutor or lecturer

- If you email us about a general inquiry, we may ask you to re-post your question in the forum
- Please include COMP90049 in email subject



I am looking for 2-3 Student Representatives

- Communication channel between class and teaching team
- Collect and pass on (anonymous) feedback or complaints
- Attend a student-staff meeting during the semester (TBD)
- Represent the **diversity** of the class

Interested? Send me an email with a short paragraph on why you want this role.



Interaction and Engagement

- We'll experiment with breakout rooms, polls, shared whiteboards... please engage!
- Feel free to ask questions / use the chat / raise your hands (I'll do my best to monitor)
- Feedback surveys
- You are encouraged to switch on your camera in lectures and (particularly) workshops to maximize engagement. Please see the recent announcement / post on the subject Home page for acknowledgment of and details on privacy concerns.

- **Topics** include: classification, clustering, optimization, unsupervised learning, semi-supervised learning, neural networks
- All from a theoretical and practical perspective
- Refreshers on maths and programming basics
- Theory in the lectures (some live-coding and demo-ing of libraries and toolkits)
- Hands-on experience in workshops and projects
- **Guest lecture 1:** academic writing skills
- **Guest lecture 2:** Industry talk with focus on bias and fairness in machine learning



Programming concepts

- We will be using **Python** and **Jupyter Notebooks**
- Basic familiarity with libraries (numpy, scikit-learn, scipy)
- You need to be able to write code to process your data, apply different algorithms, and evaluate the output
- Optional practice / demo Jupyter notebooks (most weeks)
- Optional **coding consultation sessions** in the first weeks of semester

Programming concepts

- We will be using **Python** and **Jupyter Notebooks**
- Basic familiarity with libraries (numpy, scikit-learn, scipy)
- You need to be able to write code to process your data, apply different algorithms, and evaluate the output
- Optional practice / demo Jupyter notebooks (most weeks)
- Optional **coding consultation sessions** in the first weeks of semester

Mathematical concepts

- formal maths notation
- basic probability, statistics, calculus, geometry, linear algebra
- (why?)



What Level of Maths are we Talking?

$$\ln \frac{P(y = \text{true}|x)}{1 - P(y = \text{true}|x)} = w \cdot f$$

$$\frac{P(y = \text{true}|x)}{1 - P(y = \text{true}|x)} = e^{w \cdot f}$$

$$P(y = \text{true}|x) = e^{w \cdot f} - e^{w \cdot f} P(y = \text{true}|x)$$

$$P(y = \text{true}|x) + e^{w \cdot f} P(y = \text{true}|x) = e^{w \cdot f}$$

$$P(y = \text{true}|x) = h(x) = \frac{e^{w \cdot f}}{1 + e^{w \cdot f}} = \frac{1}{1 + e^{-w \cdot f}}$$

$$P(y = \text{false}|x) = \frac{1}{1 + e^{w \cdot f}} = \frac{e^{-w \cdot f}}{1 + e^{-w \cdot f}}$$



What Level of Maths are we Talking?

$$P(y = 1|x; \beta) = h_{\beta}(x)$$

$$P(y = 0|x; \beta) = 1 - h_{\beta}(x)$$

$$\rightarrow P(y|x; \beta) = (h_{\beta}(x))^y * (1 - h_{\beta}(x))^{1-y}$$

$$\begin{aligned} & \operatorname{argmax}_{\beta} \prod_{i=1}^n P(y_i|x_i; \beta) \\ &= \operatorname{argmax}_{\beta} \prod_{i=1}^n (h_{\beta}(x_i))^{y_i} * (1 - h_{\beta}(x_i))^{1-y_i} \\ &= \operatorname{argmax}_{\beta} \sum_{i=1}^n y_i \log h_{\beta}(x_i) + (1 - y_i) \log(1 - h_{\beta}(x_i)) \end{aligned}$$



Two small coding projects (30%)

- Project 1: release week 2, due week 3
- Project 2: release week 5, due week 6
- Read in data, apply ML algorithm(s), evaluate.

Open-ended research project (30%)

- Release week 7, due week 10
- You will be given a data set and will formulate a research question and write a short research paper on your findings. You will be graded based on the quality of your report.

Final exam (40%)

- during exam period
- 2 hours; closed-book
- **Hurdle requirement:** you have to pass the exam ($\geq 50\%$).











Academic Honesty

- Videos & Quiz
- Linked from Canvas 'Home' page (or in Modules)
- CIS-specific scenarios

CIS Academic Honesty Training



Complete All Items

Videos

 Getting help from non-student friends Mark as done	<input type="radio"/>
 Copying the answer from a fellow student Mark as done	<input type="radio"/>
 Getting help from fellow students Mark as done	<input type="radio"/>
 Copying the answer from online sources	
 Do not outsource assignments Mark as done	<input type="radio"/>
 Lock screen when leaving your monitor Mark as done	<input type="radio"/>
 Protect your code (Do not share on Github) Mark as done	<input type="radio"/>
 Working and discussing with friends in the right way Mark as done	<input type="radio"/>

Quiz

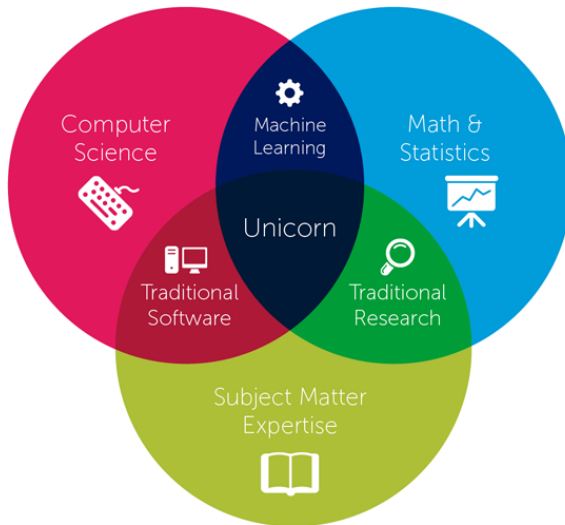
Further information

 Academic Integrity Principles at Unimelb
 Further Resources



What and Why of Machine Learning?

What is Machine Learning?




What is Machine Learning?



<https://xkcd.com/1838/>

(you're sitting in the right class!)

Top 20 Emerging Jobs



Source: <https://www.springboard.com/blog/machine-learning-engineer-salary-guide/>

Three ingredients for machine learning

... and related questions

Three ingredients for machine learning

... and related questions

1. Data

- Discrete vs continuous vs ...
- Big data vs small data
- Labeled data vs unlabeled data
- Public vs sensitive data



Three ingredients for machine learning

... and related questions

Models

- function mapping from inputs to outputs
- motivated by a data *generating* hypothesis
- probabilistic machine learning models
- geometric machine learning models
- parameters of the function are unknown



Three ingredients for machine learning

... and related questions

Learning

- Improving (on a task) after data is taken into account
- Finding the best model parameters (for a given task)
- Supervised vs. unsupervised learning



ML Example Problem

- Scenario 1

You are an archaeologist in charge of classifying a mountain of fossilized bones, and want to quickly identify any “finds of the century” before sending the bones off to a museum

- Solution:

Identify bones which are of different size/dimensions/characteristics to others in the sample and/or pre-identified bones

- Scenario 1

You are an archaeologist in charge of classifying a mountain of fossilized bones, and want to quickly identify any “finds of the century” before sending the bones off to a museum

- Solution:

Identify bones which are of different size/dimensions/characteristics to others in the sample and/or pre-identified bones

CLUSTERING/OUTLIER DETECTION

- Scenario 2:

You are an archaeologist in charge of classifying a mountain of fossilized bones, and want to come up with a consistent way of determining the species and type of each bone which doesn't require specialist skills

- Solution:

Identify some easily measurable properties of bones (size, shape, number of “lumps”, ...) and compare any new bones to a pre-classified database of bones

- Scenario 2:

You are an archaeologist in charge of classifying a mountain of fossilized bones, and want to come up with a consistent way of determining the species and type of each bone which doesn't require specialist skills

- Solution:

Identify some easily measurable properties of bones (size, shape, number of “lumps”, ...) and compare any new bones to a pre-classified database of bones

SUPERVISED CLASSIFICATION ;



- Scenario 3:

You are in charge of developing the next “release” of Coca Cola, and want to be able to estimate how well received a given recipe will be

- Solution:

Carry out taste tests over various “recipes” with varying proportions of sugar, caramel, caffeine, phosphoric acid, coca leaf extract, ... (and any number of “secret” new ingredients), and estimate the function which predicts customer satisfaction from these numbers

- Scenario 3:

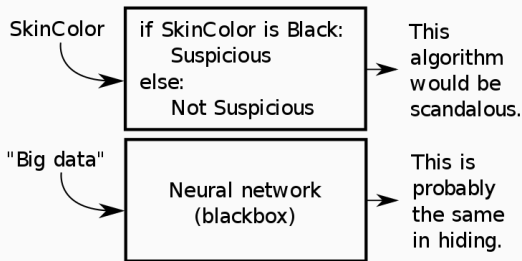
You are in charge of developing the next “release” of Coca Cola, and want to be able to estimate how well received a given recipe will be

- Solution:

Carry out taste tests over various “recipes” with varying proportions of sugar, caramel, caffeine, phosphoric acid, coca leaf extract, ... (and any number of “secret” new ingredients), and estimate the function which predicts customer satisfaction from these numbers

REGRESSION

- natural language processing
- image classification
- stock market prediction
- movie recommendation
- web search
- medical diagnoses
- spam / malware detection
- ...



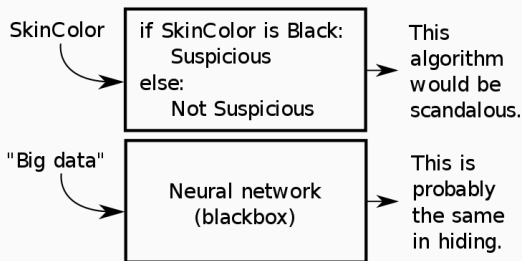
[commons.wikimedia.org/wiki/File:Pseudo-](https://commons.wikimedia.org/wiki/File:Pseudo-algorithm_comparison_for_my_slides_on_machine_learning_ethics.svg)

[algorithm_comparison_for_my_slides_on_machine_learning_ethics.svg](#)

Def 1. **Discrimination**= To make distinctions.

For example, in supervised ML, for a given instance, we might try to discriminate between the various possible classes.





[commons.wikimedia.org/wiki/File:Pseudo-](https://commons.wikimedia.org/wiki/File:Pseudo-algorithm_comparison_for_my_slides_on_machine_learning_ethics.svg)

[algorithm_comparison_for_my_slides_on_machine_learning_ethics.svg](#)

Def 2. **Discrimination**= To make decisions based on prejudice.

Digital computers have no volition, and consequently cannot be prejudiced.

However, the data may contain information which leads to an application where the ensuing behavior is prejudicial, intentionally or otherwise.



RETAIL OCTOBER 11, 2018 / 10:04 AM / UPDATED 2 YEARS AGO


Amazon scraps secret AI recruiting tool that showed bias against women

By Jeffrey Dastin

8 MIN READ



SAN FRANCISCO (Reuters) - Amazon.com Inc's [AMZN.O](#) machine-learning specialists uncovered a big problem: their new recruiting engine did not like women.



MEDIA AND TELECOMS MARCH 25, 2016 / 9:55 AM / UPDATED 5 YEARS AGO

Microsoft's AI Twitter bot goes dark after racist, sexist tweets

By Amy Tenncry, Gina Chereilus

3 MIN READ



(Reuters) - Tay, Microsoft Corp's so-called chatbot that uses artificial intelligence to engage with millennials on Twitter, lasted less than a day before it was hobbled by a barrage of racist and sexist comments by Twitter users that it parroted back to them.



Machine Learning gone wrong...

BBC Sign in Home News Sport Reel Worklife Travel

NEWS

Home | Coronavirus | Video | World | Asia | UK | Business | Tech | Science | Stories | Entertainment & Arts

Tech

Facial recognition to 'predict criminals' sparks row over AI bias

© 24 June 2020



GETTY IMAGES

A US university's claim it can use facial recognition to "predict criminality" has renewed debate over racial bias in technology.

Harrisburg University researchers said their software "can predict if someone is a criminal, based solely on a picture of their face".

Not everything that *can* be done, *should* be done

- Attributes in the data can encode information in an indirect way
- For example, home address and occupation can be used (perhaps with other seemingly-banal data) to infer age and social standing of an individual
- Potential legal exposure due to implicit “knowledge” used by a classifier
- Just because you didn’t realize doesn’t mean that you shouldn’t have realized, or at least, made reasonable efforts to check

Questions to Ask

- Who is permitted to access the data?
- For what purpose was the data collected?
- What kinds of conclusions are legitimate?
- If our conclusions defy common sense, are there confounding factors?
- Could my research / application be abused (*dual use*)?



Today

- COMP90049 Overview
- What is machine learning?
- Why is it important? Some use cases.
- What can go wrong?

Next lecture: Concepts in machine learning

Jacob Eisenstein. Natural Language Processing. MIT Press (2019)

Marc Peter Deisenroth, A Aldo Faisal, and Cheng Soon Ong. Mathematics for Machine Learning. Cambridge University Press (forthcoming)

Chris Bishop. Pattern Recognition and Machine Learning. Springer (2009)

Tom Mitchell. Machine Learning. McGraw-Hill, New York, USA (1997).



Microsoft's AI robot goes dark.

[https:](https://www.reuters.com/article/us-microsoft-twitter-bot-idUSKCN0WQ2LA)

[//www.reuters.com/article/us-microsoft-twitter-bot-idUSKCN0WQ2LA](https://www.reuters.com/article/us-microsoft-twitter-bot-idUSKCN0WQ2LA)

Amazon scraps secret recruiting tool.

[https://www.reuters.com/article/](https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G)

[us-amazon-com-jobs-automation-insight-idUSKCN1MK08G](https://www.reuters.com/article/us-amazon-com-jobs-automation-insight-idUSKCN1MK08G)

Predictive policing algorithms are biased.

<https://www.bbc.com/news/technology-53165286>