

Lecture 10: Logistic Regression

COMP90049

Introduction to Machine Learning

Semester 2, 2021

Lida Rashidi, CIS

Copyright @ University of Melbourne 2021. All rights reserved. No part of the publication may be reproduced in any form by print, photoprint, microfilm or any other means without written permission from the author.

Acknowledgement: Lea Frermann



Sofar...

- Naive Bayes + KNN
- Optimization (closed-form and iterative)
- Evaluation + Feature Selection

Today : more classification!

- Logistic Regression

Logistic Regression

Recall **Naive Bayes**

$$P(x, y) = P(y)P(x|y) = \prod_{i=1}^N P(y^i) \prod_{m=1}^M P(x_m^i | y^i)$$

- a **probabilistic generative model** of the joint probability $P(x, y)$
- optimized to maximize the likelihood of the observed data
- **naive** due to unrealistic feature independence assumptions

Recall **Naive Bayes**

$$P(x, y) = P(y)P(x|y) = \prod_{i=1}^N P(y^i) \prod_{m=1}^M P(x_m^i | y^i)$$

- a **probabilistic generative model** of the joint probability $P(x, y)$
- optimized to maximize the likelihood of the observed data
- **naive** due to unrealistic feature independence assumptions

For **prediction**, we apply **Bayes Rule** to obtain the conditional distribution

$$P(x, y) = P(y)P(x|y) = P(y|x)P(x)$$

$$P(y|x) = \frac{P(y)P(x|y)}{P(x)}$$

$$\hat{y} = \operatorname{argmax}_y P(y|x) \approx P(y)P(x|y)$$

How about we model $P(y|x)$ directly? → **Logistic Regression**



Logistic Regression on a high level

- Is a **binary** classification model
- Is a **probabilistic discriminative model** because it optimizes $P(y|x)$ directly
- Learns to optimally discriminate between inputs which belong to different classes
- No model of $P(x|y) \rightarrow$ no conditional feature independence assumption

Aside: Linear Regression

- Regression: predict a real-valued quantity y given features x , e.g.,

housing price	given	{location, size, age, ...}
success of movie (\$)	given	{cast, genre, budget, ...}
air quality	given	{temperature, timeOfDay, CO2, ...}



Aside: Linear Regression

- Regression: predict a real-valued quantity y given features x , e.g.,

housing price	given	{location, size, age, ...}
success of movie (\$)	given	{cast, genre, budget, ...}
air quality	given	{temperature, timeOfDay, CO2, ...}

- linear regression** is the simplest regression model
- a real-valued \hat{y} is predicted as a linear combination of weighted feature values

$$\begin{aligned}\hat{y} &= \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots \\ &= \theta_0 + \sum_i \theta_i x_i\end{aligned}$$

- The weights $\theta_0, \theta_1, \dots$ are model parameters, and need to be optimized during training
- Loss (error) is the sum of squared errors (SSE): $L = \sum_{i=1}^N (\hat{y}^i - y^i)^2$



Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is `true` (1) or `false` (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach
- How about: $p(x)$ as a linear function of x . Problem: probabilities are bounded in 0 and 1, linear functions are not.

$$p(x) = \theta_0 + \theta_1 x_1 + \dots \theta_F x_F$$

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach
- ~~How about: $p(x)$ as a linear function of x . Problem: probabilities are bounded in 0 and 1, linear functions are not.~~

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach
- ~~How about: $p(x)$ as a linear function of x . Problem: probabilities are bounded in 0 and 1, linear functions are not.~~
- How about: $\log p(x)$ as a linear function of x . Problem: \log is bounded in one direction, linear functions are not.

$$\log p(x) = \theta_0 + \theta_1 x_1 + \dots \theta_F x_F$$

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach
- ~~How about: $p(x)$ as a linear function of x . Problem: probabilities are bounded in 0 and 1, linear functions are not.~~
- ~~How about: $\log p(x)$ as a linear function of x . Problem: \log is bounded in one direction, linear functions are not.~~

Logistic Regression: Derivation I

- Let's assume a **binary** classification task, y is true (1) or false (0).
- We model **probabilities** $P(y = 1|x; \theta) = p(x)$ as a function of observations x under parameters θ . [What about $P(y = 0|x; \theta)$?]
- We want to use a **regression** approach
- ~~How about: $p(x)$ as a linear function of x . Problem: probabilities are bounded in 0 and 1, linear functions are not.~~
- ~~How about: $\log p(x)$ as a linear function of x . Problem: \log is bounded in one direction, linear functions are not.~~
- How about: minimally modifying $\log p(x)$ such that is is unbounded, by applying the **logistic** transformation

$$\log \frac{p(x)}{1 - p(x)} = \theta_0 + \theta_1 x_1 + \dots \theta_F x_F$$



$$\log \frac{p(x)}{1 - p(x)} = \theta_0 + \theta_1 x_1 + \dots \theta_F x_F$$

- also called the **log odds**
- the **odds** are defined as the fraction of success over the fraction of failures

$$\text{odds} = \frac{P(\text{success})}{P(\text{failures})} = \frac{P(\text{success})}{1 - P(\text{success})}$$

- e.g., the odds of rolling a 6 with a fair dice are:

$$\frac{1/6}{1 - (1/6)} = \frac{0.17}{0.83} = 0.2$$


Logistic Regression: Derivation III

$$\log \frac{P(x)}{1 - P(x)} = \theta_0 + \theta_1 x_1 + \dots \theta_F x_F$$

If we rearrange and solve for $P(x)$, we get

$$\begin{aligned} P(x) &= \frac{\exp(\theta_0 + \theta_1 x_1 + \dots \theta_F x_F)}{1 + \exp(\theta_0 + \theta_1 x_1 + \dots \theta_F x_F)} = \frac{\exp(\theta_0 + \sum_{f=1}^F \theta_f x_f)}{1 + \exp(\theta_0 + \sum_{f=1}^F \theta_f x_f)} \\ &= \frac{1}{1 + \exp(-(\theta_0 + \theta_1 x_1 + \dots \theta_F x_F))} = \frac{1}{1 + \exp(-(\theta_0 + \sum_{f=1}^F \theta_f x_f))} \end{aligned}$$


- where the RHS is the inverse logit (or **logistic function**)
- we pass a regression model through the logistic function to obtain a valid probability prediction



$$P(y|x; \theta) = \frac{1}{1 + \exp(-(\theta_0 + \sum_{f=1}^F \theta_f x_f))}$$

A closer look at the logistic function

Most inputs lead to $P(y|x)=0$ or $P(y|x)=1$. That is intended, because all true labels are either 0 or 1.




- $(\theta_0 + \sum_{f=1}^F \theta_f x_f) > 0$ means $y = 1$
- $(\theta_0 + \sum_{f=1}^F \theta_f x_f) \approx 0$ means most uncertainty
- $(\theta_0 + \sum_{f=1}^F \theta_f x_f) < 0$ means $y = 0$

Logistic Regression: Prediction

- The logistic function returns the probability of $P(y = 1)$ given an input x

$$P(y = 1 | x_1, x_2, \dots, x_F; \theta) = \frac{1}{1 + \exp(-(\theta_0 + \sum_{f=1}^F \theta_f x_f))} = \sigma(x; \theta)$$

- We define a **decision boundary**, e.g., predict $y = 1$ if $P(y = 1 | x_1, x_2, \dots, x_F; \theta) > 0.5$ and $y = 0$ otherwise



Example!

$$P(y = 1|x_1, x_2, \dots, x_F; \theta) = \frac{1}{1 + \exp(-(\theta_0 + \sum_{f=1}^F \theta_f x_f))} = \frac{1}{1 + \exp(-(\theta^T x))} = \sigma(\theta^T x)$$

Model parameters

$$\theta = [0.1, -3.5, 0.7, 2.1]$$

(Small) Test Data set

Outlook	Temp	Humidity	Class
<i>rainy</i>	<i>cool</i>	<i>normal</i>	0
<i>sunny</i>	<i>hot</i>	<i>high</i>	1

Feature Function

$$x_0 = 1 \text{ (bias term)}$$

$$x_1 = \begin{cases} 1 & \text{if outlook=sunny} \\ 2 & \text{if outlook=overcast} \\ 3 & \text{if outlook=rainy} \end{cases}$$

$$x_2 = \begin{cases} 1 & \text{if temp=hot} \\ 2 & \text{if temp=mild} \\ 3 & \text{if temp=cool} \end{cases}$$

$$x_3 = \begin{cases} 1 & \text{if humidity=normal} \\ 2 & \text{if humidity=high} \end{cases}$$



Example!

$$P(y = 1|x_1, x_2, \dots, x_F; \theta) = \frac{1}{1 + \exp(-(\theta_0 + \sum_{f=1}^F \theta_f x_f))} = \frac{1}{1 + \exp(-(\theta^T x))} = \sigma(\theta^T x)$$

Model parameters

$$\theta = [0.1, -3.5, 0.7, 2.1]$$

(Small) Test Data set

Outlook	Temp	Humidity	Class
<i>rainy</i>	<i>cool</i>	<i>normal</i>	0
<i>sunny</i>	<i>hot</i>	<i>high</i>	1

Feature Function

$$x_0 = 1 \text{ (bias term)}$$

$$x_1 = \begin{cases} 1 & \text{if outlook=sunny} \\ 2 & \text{if outlook=overcast} \\ 3 & \text{if outlook=rainy} \end{cases}$$

$$x_2 = \begin{cases} 1 & \text{if temp=hot} \\ 2 & \text{if temp=mild} \\ 3 & \text{if temp=cool} \end{cases}$$

$$x_3 = \begin{cases} 1 & \text{if humidity=normal} \\ 2 & \text{if humidity=high} \end{cases}$$



What are the four steps we would follow in finding the optimal parameters?

Mimimize the Negative conditional log likelihood

$$\mathcal{L}(\theta) = -P(Y|X; \theta) = -\prod_{i=1}^N P(y^i|x^i; \theta)$$

note that

$$P(y = 1|x; \theta) = \sigma(\theta^T x)$$

$$P(y = 0|x; \theta) = 1 - \sigma(\theta^T x)$$

Mimimize the **Negative conditional log likelihood**

$$\mathcal{L}(\theta) = -P(Y|X; \theta) = -\prod_{i=1}^N P(y^i|x^i; \theta)$$

note that

$$P(y = 1|x; \theta) = \sigma(\theta^T x)$$

$$P(y = 0|x; \theta) = 1 - \sigma(\theta^T x)$$

so

$$\begin{aligned}\mathcal{L}(\theta) &= -P(Y|X; \theta) = -\prod_{i=1}^N P(y^i|x^i; \theta) \\ &= -\prod_{i=1}^N (\sigma(\theta^T x^i))^{y^i} * (1 - \sigma(\theta^T x^i))^{1-y^i}\end{aligned}$$

Mimimize the Negative conditional log likelihood

$$\mathcal{L}(\theta) = -P(Y|X; \theta) = -\prod_{i=1}^N P(y^i|x^i; \theta)$$

note that

$$P(y = 1|x; \theta) = \sigma(\theta^T x)$$

$$P(y = 0|x; \theta) = 1 - \sigma(\theta^T x)$$

so

$$\begin{aligned}\mathcal{L}(\theta) &= -P(Y|X; \theta) = -\prod_{i=1}^N P(y^i|x^i; \theta) \\ &= -\prod_{i=1}^N (\sigma(\theta^T x^i))^{y^i} * (1 - \sigma(\theta^T x^i))^{1-y^i}\end{aligned}$$

take the log of this function

$$\log \mathcal{L}(\theta) = -\sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$



Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

Also

- Derivative of sum = sum of derivatives \rightarrow focus on 1 training input
- Compute $\frac{\partial \mathcal{L}}{\partial \theta_j}$ for each θ_j individually, so focus on 1 θ_j



Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$



Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$

$$\begin{array}{c} \downarrow \\ \frac{\partial \log \mathcal{L}(\theta)}{\partial p} = -\left(\frac{y}{p} - \frac{1-y}{1-p} \right) \end{array}$$

$$(\text{ because } \mathcal{L}(\theta) = -[y \log p + (1-y) \log(1-p)]$$




Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$


$$\frac{\partial p}{\partial z} = \frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$$




Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$


$$\frac{\partial z}{\partial \theta_j} = \frac{\partial \theta^T x}{\partial \theta_j} = x_j$$

Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\begin{aligned} \frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} &= \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x \\ &= -\frac{y}{p} - \frac{1-y}{1-p} \times \sigma(z)[1 - \sigma(z)] \times x_j \end{aligned}$$



Take 1st Derivative of the Objective Function

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial D} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C} \times \frac{\partial C}{\partial D}$

$$\begin{aligned} \frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} &= \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x \\ &= -\frac{y}{p} - \frac{1-y}{1-p} \times \sigma(z)[1 - \sigma(z)] \times x_j \\ &= [\sigma(\theta^T x) - y] \times x_j \end{aligned}$$



Logistic Regression: Parameter Estimation III

The derivative of the log likelihood wrt. a single parameter θ_j for **all** training examples

$$\frac{\log \mathcal{L}(\theta)}{\partial \theta_j} = \sum_{i=1}^N \left(\sigma(\theta^T x^i) - y^i \right) x_j^i$$

- Now, we would set derivatives to zero (**Step 3**) and solve for θ (**Step 4**)
- Unfortunately, that's not straightforward here (as for Naive Bayes)
- Instead, we will use an iterative method: **Gradient Descent**



Logistic Regression: Parameter Estimation III

The derivative of the log likelihood wrt. a single parameter θ_j for **all** training examples

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \sum_{i=1}^N \left(\sigma(\theta^T x^i) - y^i \right) x_j^i$$

- Now, we would set derivatives to zero (**Step 3**) and solve for θ (**Step 4**)
- Unfortunately, that's not straightforward here (as for Naive Bayes)
- Instead, we will use an iterative method: **Gradient Descent**

$$\theta_j^{(new)} \leftarrow \theta_j^{(old)} - \eta \frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j}$$

$$\theta_j^{(new)} \leftarrow \theta_j^{(old)} - \eta \sum_{i=1}^N \left(\sigma(\theta^T x^i) - y^i \right) x_j^i$$



Multinomial Logistic Regression

- So far we looked at problems where either $y = 0$ or $y = 1$ (e.g., spam classification: $y \in \{\text{play}, \text{not_play}\}$)

$$P(y = 1|x; \theta) = \sigma(\theta^T x) = \frac{\exp(\theta^T x)}{1 + \exp(\theta^T x)}$$

$$P(y = 0|x; \theta) = 1 - \sigma(\theta^T x) = 1 - \frac{\exp(\theta^T x)}{1 + \exp(\theta^T x)}$$



Multinomial Logistic Regression

- So far we looked at problems where either $y = 0$ or $y = 1$ (e.g., spam classification: $y \in \{\text{play}, \text{not_play}\}$)

$$P(y = 1|x; \theta) = \sigma(\theta^T x) = \frac{\exp(\theta^T x)}{1 + \exp(\theta^T x)}$$

$$P(y = 0|x; \theta) = 1 - \sigma(\theta^T x) = 1 - \frac{\exp(\theta^T x)}{1 + \exp(\theta^T x)}$$

- But what if we have more than 2 classes, e.g., $y \in \{\text{positive}, \text{negative}, \text{neutral}\}$
- we predict the probability of each class c by passing the input representation through the **softmax** function, a generalization of the sigmoid

$$p(y = c|x; \theta) = \frac{\exp(\theta_c x)}{\sum_k \exp(\theta_k x)}$$

- we learn a parameter vector θ_c for each class c



Example! Multi-class with 1-hot features

$$p(y = c|x; \theta) = \frac{\exp(\theta_c x)}{\sum_k \exp(\theta_k x)}$$

Model parameters

$$\theta_{c0} = [0.1, -3.5, 0.7, 2.1]$$

$$\theta_{c1} = [0.6, 2.5, 2.7, -2.1]$$

$$\theta_{c2} = [3.1, 1.5, 0.07, 3.6]$$

(Small) Test Data set

Outlook	Temp	Humidity	Class
<i>rainy</i>	<i>cool</i>	<i>normal</i>	0 (don't play)
<i>sunny</i>	<i>cool</i>	<i>normal</i>	1 (maybe play)
<i>sunny</i>	<i>hot</i>	<i>high</i>	2 (play)

Feature Function

$$x_0 = 1 \text{ (bias term)}$$

$$x_1 = \begin{cases} 1 & \text{if outlook=sunny} \\ 2 & \text{if outlook=overcast} \\ 3 & \text{if outlook=rainy} \end{cases}$$

$$x_2 = \begin{cases} 1 & \text{if temp=hot} \\ 2 & \text{if temp=mild} \\ 3 & \text{if temp=cool} \end{cases}$$

$$x_3 = \begin{cases} 1 & \text{if humidity=normal} \\ 2 & \text{if humidity=high} \end{cases}$$

$$x_0 = 1 \text{ (bias term)}$$

$$x_1 = \begin{cases} [100] & \text{if outlook=sunny} \\ [010] & \text{if outlook=overcast} \\ [001] & \text{if outlook=rainy} \end{cases}$$

$$x_2 = \begin{cases} [100] & \text{if temp=hot} \\ [010] & \text{if temp=mild} \\ [001] & \text{if temp=cool} \end{cases}$$

$$x_3 = \begin{cases} [10] & \text{if humidity=normal} \\ [01] & \text{if humidity=high} \end{cases}$$



Example! Multi-class with 1-hot features

$$p(y = c|x; \theta) = \frac{\exp(\theta_c x)}{\sum_k \exp(\theta_k x)}$$

(Small) Test Data set

Outlook	Temp	Humidity	Class
0 0 1	0 0 1	1 0	0
1 0 0	0 0 1	1 0	1
1 0 0	1 0 0	0 1	2

Feature Function

$x_0 = 1$ (bias term)

$x_1 = \begin{cases} [100] & \text{if outlook=sunny} \\ [010] & \text{if outlook=overcast} \\ [001] & \text{if outlook=rainy} \end{cases}$

$x_2 = \begin{cases} [100] & \text{if temp=hot} \\ [010] & \text{if temp=mild} \\ [001] & \text{if temp=cool} \end{cases}$

$x_3 = \begin{cases} [10] & \text{if humidity=normal} \\ [01] & \text{if humidity=high} \end{cases}$



Pros

- Probabilistic interpretation
- No restrictive assumptions on features
- Often outperforms Naive Bayes
- Particularly suited to frequency-based features (so, popular in NLP)

Cons

- Can only learn *linear* feature-data relationships
- Some feature scaling issues
- Often needs a lot of data to work well
- Regularisation a nuisance, but important since overfitting can be a big problem



- Derivation of logistic regression
- Prediction
- Derivation of maximum likelihood

References

Cosma Shalizi. *Advanced Data Analysis from an Elementary Point of View*. Chapters 11.1 and 11.2. Online Draft.

<https://www.stat.cmu.edu/~cshalizi/ADAfaEPoV/ADAfaEPoV.pdf>

Dan Jurafsky and James H. Martin. *Speech and Language Processing*. Chapter 5. Online Draft V3.0.

<https://web.stanford.edu/~jurafsky/slp3/>



Optional: Detailed Parameter Estimation

Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$



Optional: Detailed Parameter Estimation

Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

Also

- Derivative of sum = sum of derivatives \rightarrow focus on 1 training input
- Compute $\frac{\partial \mathcal{L}}{\partial \theta_j}$ for each θ_j individually, so focus on 1 θ_j



Optional: Detailed Parameter Estimation

Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$



Optional: Detailed Parameter Estimation


Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$


$$\frac{\partial \log \mathcal{L}(\theta)}{\partial p} = -\left(\frac{y}{p} - \frac{1-y}{1-p} \right)$$

(because $\mathcal{L}(\theta) = -[y \log p + (1 - y) \log(1 - p)]$)



Optional: Detailed Parameter Estimation


Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$


$$\frac{\partial p}{\partial z} = \frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$$



Optional: Detailed Parameter Estimation


Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$


$$\frac{\partial z}{\partial \theta_j} = \frac{\partial \theta^T x}{\partial \theta_j} = x_j$$



Optional: Detailed Parameter Estimation

Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$

$$= - \left[\frac{y}{p} - \frac{1-y}{1-p} \right] \times \sigma(z)[1 - \sigma(z)] \times x_j \quad \text{[[combine 3 derivatives]]}$$

$$= - \left[\frac{y}{p} - \frac{1-y}{1-p} \right] \times p[1 - p] \times x_j \quad \text{[[} \sigma(z) = p \text{]]}$$

$$= - \left[\frac{y(1-p)}{p(1-p)} - \frac{p(1-y)}{p(1-p)} \right] \times p[1 - p] \times x_j \quad \text{[[} \times \frac{1-p}{1-p} \text{ and } \frac{p}{p} \text{]]}$$

$$= - [y(1-p) - p(1-y)] \times x_j \quad \text{[[cancel terms]]}$$



Optional: Detailed Parameter Estimation

Step 2 Differentiate the loglikelihood wrt. the parameters

$$\log \mathcal{L}(\theta) = - \sum_{i=1}^N y^i \log \sigma(\theta^T x^i) + (1 - y^i) \log(1 - \sigma(\theta^T x^i))$$

Preliminaries

- The derivative of the logistic (sigmoid) function is $\frac{\partial \sigma(z)}{\partial z} = \sigma(z)[1 - \sigma(z)]$
- The chain rule tells us that $\frac{\partial A}{\partial C} = \frac{\partial A}{\partial B} \times \frac{\partial B}{\partial C}$

$$\frac{\partial \log \mathcal{L}(\theta)}{\partial \theta_j} = \frac{\partial \log \mathcal{L}(\theta)}{\partial p} \times \frac{\partial p}{\partial z} \times \frac{\partial z}{\partial \theta_j} \quad \text{where } p = \sigma(\theta^T x) \text{ and } z = \theta^T x$$

$$= - \left[y(1 - p) - p(1 - y) \right] \times x_j \quad \text{[[copy from last slide]]}$$

$$= - \left[y - yp - p + yp \right] \times x_j \quad \text{[[multiply out]]}$$

$$= - \left[y - p \right] \times x_j \quad \text{[[-yp+yp=0]]}$$

$$= \left[p - y \right] \times x_j \quad \text{[[-[y-p] = -y+p = p-y]]}$$

$$= \left[\sigma(\theta^T x) - y \right] \times x_j \quad \text{[[} p = \sigma(z), z = \theta^T x \text{]]}$$

