

Practice exam (long answer)

Edit View Insert Format Tools Table

Started: Nov 7 at 21:16

12pt ∨ Paragraph ∨ | **B** *I* U A ∨ ✎ ∨ T² ∨ |

Quiz Instructions

These are practice 'long answer' questions, just based on the sample exam questions compiled from previous years.

The subject staff have worked hard on the real exam to try to make answering questions as straightforward as possible, so hopefully it will be smoother than these sample questions, but these questions are illustrative of the *type* of long answer questions on the final exam.

Blind and heuristic search

Question 1

0 pts

The following search algorithms can be implemented in a similar manner, differing only in the abstract data structure used to implement the open list.

- Breadth First Search
- Depth First Search
- Uniform Cost Search






Answer the following questions:

- I. State which abstract data structure you would use for Breadth First Search. Briefly explain, in one or two sentences, why you would choose the particular data structure.
- II. State which abstract data structure you would use for Depth First Search. Briefly explain, in one or two sentences, why you would choose the particular data structure.
- III. State which abstract data structure you would use for Uniform Cost Search. Briefly explain, in one or two sentences, why you would choose the particular data structure.

Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ | **B** *I* U A ▾  ▾ T^2 ▾ |

p






  | 0 words |   **Question 2****0 pts**

Recall that A* requires a consistent heuristic to guarantee finding an optimal plan without reopening nodes. Draw or define a graph and make up an admissible but inconsistent heuristic function where A* returns a suboptimal solution.

Write down the steps involved in the A* algorithm by devising a simple example to illustrate.

Note: avoid making large examples, a graph with 4 nodes should be sufficient.

p

  | 0 words |   

The next five questions refer to the following problem where a robot can move horizontally and vertically among adjacent cells, such that the h_{add} heuristic estimates the same values as the *Manhattan* heuristic. Note that the robot *cannot* move into the shaded area, which is inaccessible to the robot.

Edit View Insert Format Tools Table

12pt

Paragraph

B

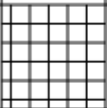
I

U

A

U

T^2

A	B	C	D
E	F	G	H
I	J		K

In answering the sub-questions below, you are allowed to use variables as arguments for the actions (action schemes), specifying the values of the variables. Note: it is **not** compulsory to use PDDL syntax, as long as you can convey the main ideas.

Hint: consider that the position of the robot could be modelled as either:

- row/column tuples, e.g. $\langle 0, 0 \rangle$ could refer to the lower left cell, or
- single cells, e.g. position $\langle I \rangle$ could refer to the lower left cell.

Question 3

0 pts

Describe briefly in STRIPS how to model the problem where a robot can move horizontally and vertically among adjacent cells, such that the h_{add} heuristic estimates the same values as the *Manhattan* heuristic

Edit View Insert Format Tools Table

12pt ▾

Paragraph ▾

B*I*UA ▾ ▾ T^2 ▾

p



0 words

**Question 4****0 pts**

Describe briefly in STRIPS how to model the problem where a robot can move horizontally and vertically among adjacent cells, such that the h_{add} heuristic estimates the same values as the *Optimal* heuristic.

p



0 words

**Question 5****0 pts**

Using your last STRIPS encoding where $h_{add} = h^*$, an initial state $s_0 = \text{robot at location } I$ and a goal state $s_g = \text{robot at location } K$, compute $h_{add}(s_g)$ and h_{ff} from the best supporters induced by h_{add} .

Show your working

Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ | **B** *I* U A ▾  ▾ T^2 ▾ |

p



0 words

**Question 6****0 pts**

We have 2 rooms A and B, 2 objects o1 and o2. A robot can load and unload objects if the robot is at the same location, and move from one room to the other. We want to get o1 and o2 into room B, given that both are initially at A.

Model the problem in STRIPS in such a way that the optimal plan would be the following:

```
pick(o1, A), move(A, B), drop(o1, B), move(B, A), pick(o2, A), move(A, B), drop(o2, B)
```

Edit View Insert Format Tools Table

12pt ∨

Paragraph ∨

B*I*UA ∨ ∨ T^2 ∨

p



0 words

**Question 7****0 pts**

What's the $h_{max}(s_0)$ value of your STRIPS model, where s_0 stands for the initial state?

Show your working

MDPs and reinforcement learning

The next six questions refer to the following example

Imagine a kitchen robot whose task is to ask the user for the cereal they want for breakfast, wait for the user answer, and then hand out the desired cereal. We want to design a simple MDP for this task.

Edit View Insert Format Tools Table

12pt ▾

Paragraph ▾

B

I

U

A ▾

 ▾

T^2 ▾

A simple way to model it is via a discount-reward MDP with only two states: *UnknownCereal*, where the robot does not know which cereal to give, and *KnownCereal*, where the robot knows the cereal to hand out.

There are only two possible actions in the model:

- Action *AskCerealType* corresponds to the robot asking the user for the cereal box he/she wishes to have. The action is only available in state *UnknownCereal*, and has a reward of -1 (a cost) in that state.

The probability of reaching state *KnownCereal* is 0.8 (if the user answers the robot's question), and otherwise (probability 0.2) the robot remains in *UnknownCereal* (if the user ignores the question or provides an unclear answer).

- Action *GiveCereal* corresponds to the robot physically giving the cereal to the user. The action is available in the state *KnownCereal*, and has a reward of 5.

When the robot executes this action in state *KnownCereal*, the MDP reaches an absorbing goal state and finishes. As such, there is no actions or reward available from this absorbing goal state.

When the robot executes this action in the *UnknownCereal* state, the MDP reaches the absorbing goal state with probability 0.3 (it gets lucky and hands the right cereal) and receives a reward of 5, or the person rejects the cereal and the robot goes back to the *UnknownCereal* state with probability 0.7 and receives a reward of -2.

Question 8

0 pts

Assuming a discount factor $\gamma = 0.9$, calculate the the value function V for each of the states *UnknownCereal* and *KnownCereal* using value iteration, for the 2nd and 3rd iterations.

Show your working.

Iteration 1:

$V(\text{KnownCereal}) = 0.0$

$V(\text{UnknownCereal}) = 0.0$

Edit View Insert Format Tools Table

12pt ∨ Paragraph ∨ | **B** *I* U A ∨ ✎ ∨ T^2 ∨ |**Question 9****0 pts**

Given the value function that you calculate above, what policy would maximise the robot's expected reward?

Show your working.

Question 10**0 pts**

In your own words, explain the difference between value iteration and policy iteration.

Edit View Insert Format Tools Table

12pt ∨ Paragraph ∨ | **B** *I* U A ∨  ∨ T² ∨ |**Question 11****0 pts**

The robot designers have found that the probabilities used for outcomes are incorrect and different for each household. As such, they decide to instead use reinforcement learning to learn the policy after deployment.

A few weeks after deployment, one such robot has the following Q-table:

State	AskCereal	GiveSerial
<i>UnknownCereal</i>	7.2	1.9
<i>KnownCereal</i>	--	3.4

Assuming a discount reward factor of 0.9 and a learning rate of 0.5.

In the state *UnknownCereal*, the robot executes *GiveCereal*, which is rejected. It decides to execute *GiveCereal* again.

Update the Q-table both the 1-step Q-learning and 1-step SARSA updates for the first *GiveCereal* action. That is, calculate two new Q-tables.

Show your working.

Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ | **B** *I* U A ▾  ▾ T^2 ▾ |

p



0 words

**Question 12****0 pts**

Challenge Question: In the state *UnknownCereal*, the robot executes *GiveCereal* three times in a row. Each time the cereal is rejected.

Calculate the 2-step SARSA update for this execution of three actions

Assume the same parameters (discounter reward factor, etc) as in the previous question

Show your working

Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ | **B** *I* U A ▾  ▾ T^2 ▾ |  ▾  ▾  ▾ | ⋮

p

Edit View Insert Format Tools Table

12pt ∨

Paragraph ∨

B*I*UA ∨ ∨T² ∨**Question 13****0 pts**

The robot designers are finding that the owners of the robots are demanding a refund because the robots get their cereal preference wrong too often. Give one technique that the robot designers could do to improve the situation.

Question 14**0 pts**

Imagine a reinforcement learning algorithm that monitors heart beat information from a fitness device, such as a FitBit, to determine whether a person develops a heart problem, such as an irregular heartbeat or a faster heartbeat.

List one potential ethical dilemma that could occur in such a situation. Justify why you believe this could be a serious problem.

Edit View Insert Format Tools Table

12pt ∨ Paragraph ∨ | **B** *I* U A ∨  ∨ T² ∨ |

Game theory

Question 15

0 pts

Consider the following abstract two-player game in normal form. Find all pure and mixed-strategy equilibria for this game.

		Player 2		
		L	M	R
Player 1	U	2,4	3,0	1,1
	D	3,2	10,3	0,4

HINT: Consider the notion of *dominated strategies*, in which some strategies are strictly dominated by others, so can be discarded.

Show your working

Edit View Insert Format Tools Table

12pt Paragraph **B** *I* U A  T^2 **Question 16****0 pts**

Challenge Question: Two players, A and B play the following game.

First A must choose IN or OUT.

If A chooses OUT the game ends, and the payoffs: are A gets 2, and B gets 0.

If A chooses IN then B observes this and must then choose IN or OUT.

If B chooses OUT the game ends, and the payoffs are: B gets 2, and A gets 0.

If B chooses IN then they play the following simultaneous move game:

		Player B	
		L	R
Player A	U	3,1	0,2
	D	-1,2	1,3

Draw the extended-form tree for this game and calculate the equilibria of the extended-form game.


Note: to 'draw' the tree, you can use text format; e.g.

U -1.a-> X

X -2.b-> Y (0,1)

X -2.c-> Z (1,0)

represents a tree with root node U. From U, L chooses between moves *b* and *c*

Alternatively, you may draw it by hand on using the  icon at the top of the text box

Edit View Insert Format Tools Table

12pt ▾

Paragraph ▾

B

I

U

A ▾

 ▾

T^2 ▾

Question 17

0 pts

Challenge Question. *The following question is under-specified, but is intended to improve your understanding of subject content and general problem-solving ability, rather than act as an entirely accurate reflection of an exam question.*

Consider a person who is mugged by someone on the street with a gun. The person has an unloaded gun in their own pocket. They could get the gun out to try to scare off the mugger, however, they risk the mugger shooting them instead. If they do not get out the unloaded gun, their mugger has time to search only their left or right pocket, but not both and the mugger will not shoot them. The person has \$100 in their left pocket and nothing in their right pocket, and the mugger knows this. Assume that if the mugger does not get the \$100, they get no payoff.

Should the person draw their unloaded gun or not?

Justify your answer

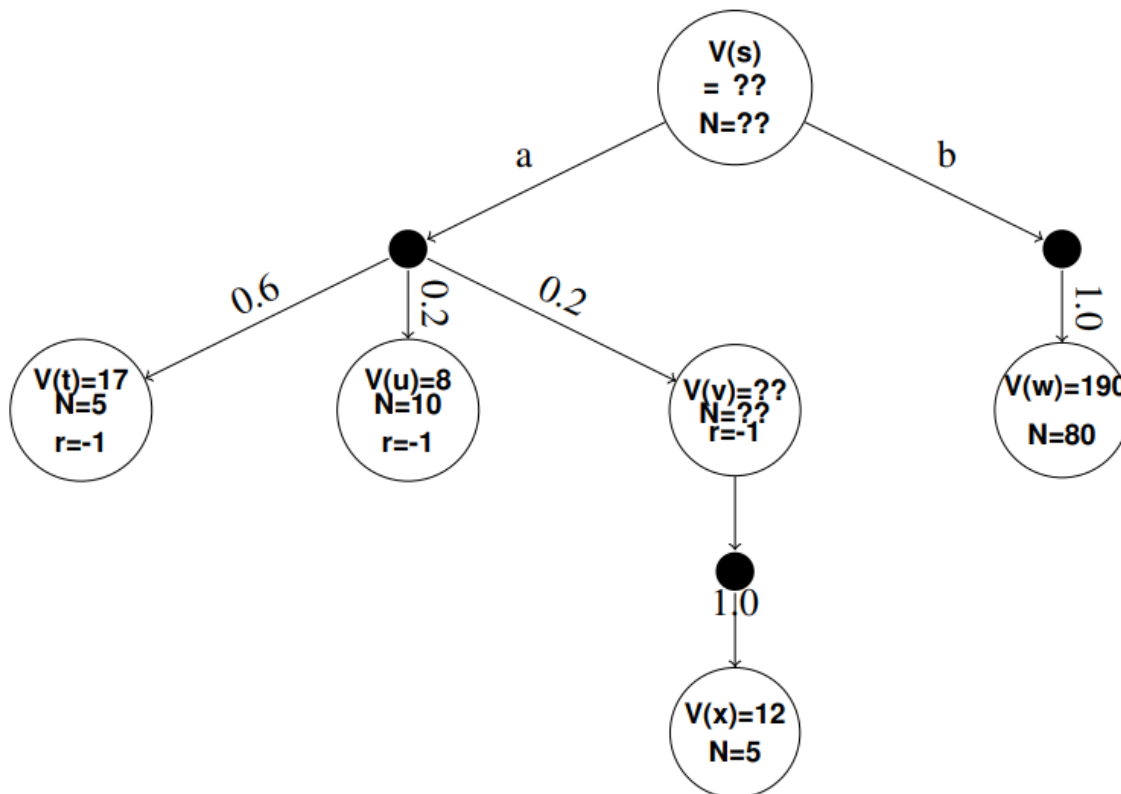
Edit View Insert Format Tools Table

12pt Paragraph **B** *I* U A ✓ ✓ T^2 ✓

Question 18

0 pts

Below is a tree from MCTS in which a collection of roll-outs have been performed, and six nodes expanded are shown (others are omitted as they are not relevant to the question). The notation $V(s) = v$ indicates that v is the value of state s , $N = n$ indicates that the node has been visited n times, and $r = R$ indicates that the reward R is given when this state is visited.



Using back-propagation, calculate the values for the states s and v .

Show your working

Edit View Insert Format Tools Table

12pt ∨

Paragraph ∨

B*I*UA ∨ ∨ T^2 ∨**Question 19****0 pts**

Given your answer for the previous question, what would action should be chosen at state s to maximise expected reward?

Show your working.

Edit View Insert Format Tools Table

12pt ▾ Paragraph ▾ | **B** *I* U A ▾  ▾ T² ▾ |

Not saved

Submit Quiz