

Dr Renata Borovica-Gajic
David Eccles



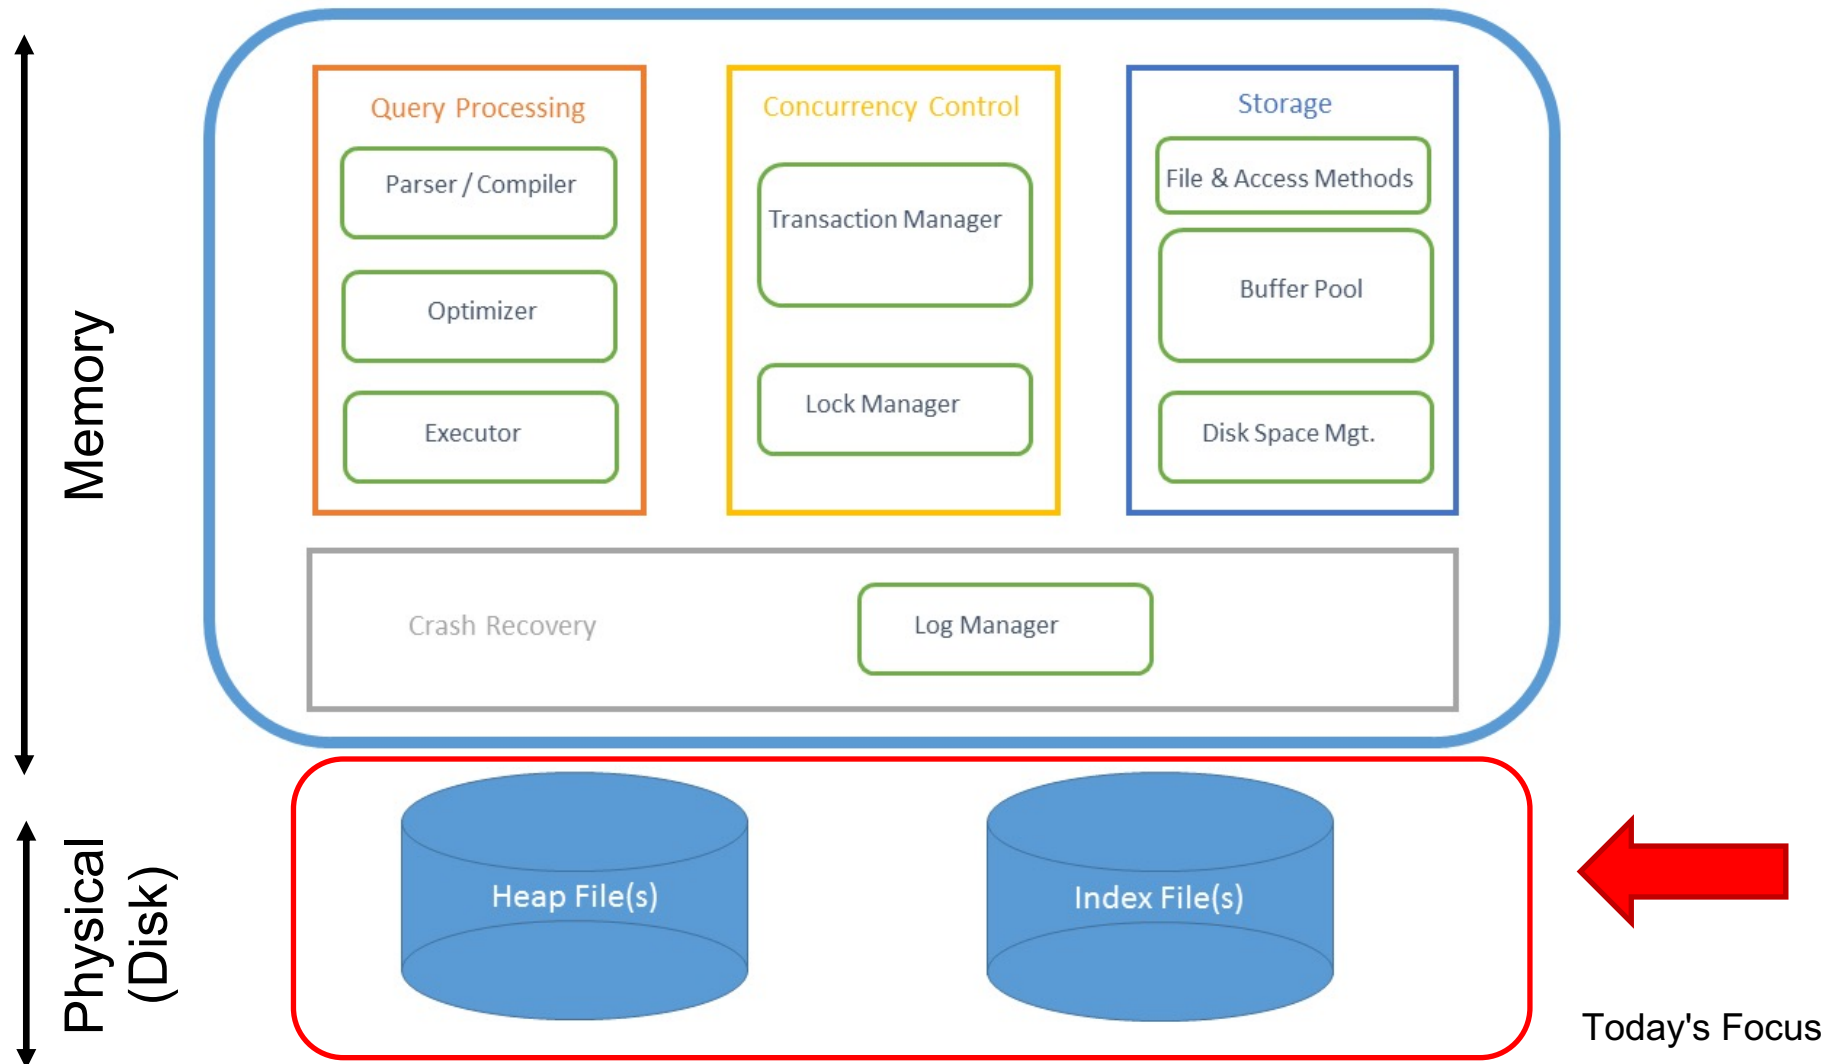
THE UNIVERSITY OF
MELBOURNE

INFO90002 Database Systems & Information Modelling

Lecture 14
Storage and Indexing



INFO90002 Database Systems & Information Modelling





- File organization (Heap & sorted files)
- Index files & indexes
- Index classification

Readings: Chapter 8, Ramakrishnan & Gehrke, Database Systems




THE UNIVERSITY OF
MELBOURNE

File Types

Heap files and Sorted files

- **FILE:** A collection of **pages**, each containing a collection of **records**.



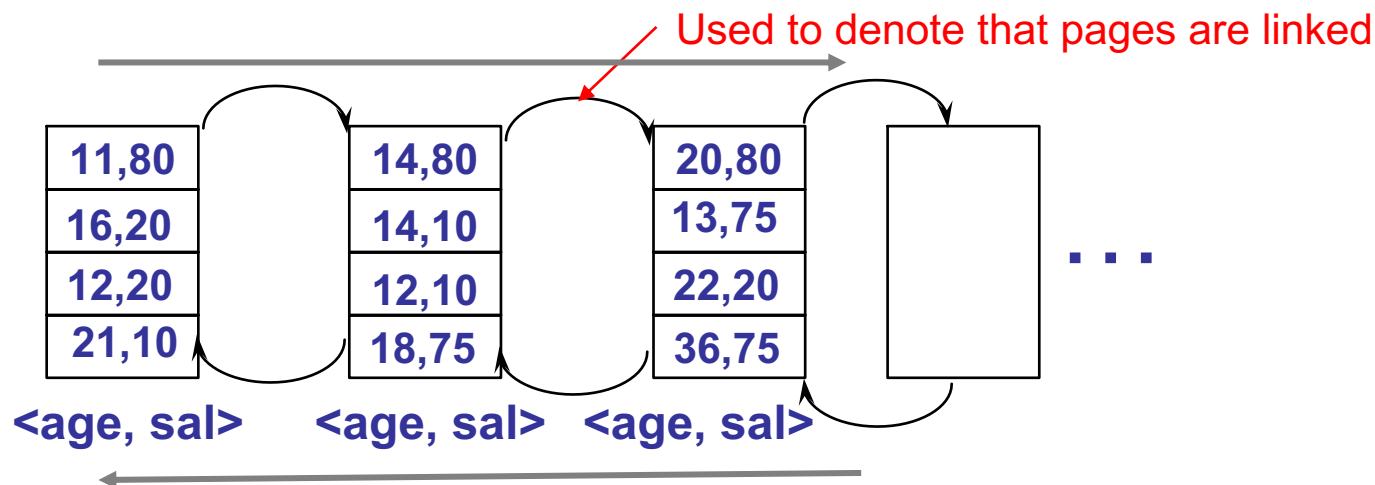
- DBMS must support:
 - insert/delete/modify record
 - read a particular record (specified using *record id*)
 - scan all records (possibly with some conditions on the records to be retrieved)

Access Methods


- Many alternatives exist, *each good for some situations, and not so good in others:*
 1. **Heap files:** no particular order among records
 - Suitable when typical access is a file scan retrieving **all** records
 2. **Sorted Files:** pages and records within pages are ordered by some condition
 - Best for retrieval (of a range of records) in some order
 3. **Index File Organizations:**
 - Special data structure that has the fastest retrieval in some order

1. Heap (Unordered) Files

- Simplest file structure, contains records in **no** particular order
- As file grows and shrinks, disk pages are allocated and de-allocated
 - Fastest for inserts compared to other alternatives



- Similar structure like heap files (pages and records), but pages and records are **ordered**
- Fast for range queries, but hard for maintenance (each insert potentially reshuffles records)
- **Example:** A sorted file ordered by *age*



INFO90002

- DBMS model the **cost** of all operations
- The cost is typically expressed in **the number of page accesses** (or disk I/O operations)
 - 1 page access (on disk) == 1 I/O (used interchangeably)
- **Example:** If we have a table of 100 records, and each page can store 10 records, what would be the cost of accessing the entire file
- **Answer:** For 100 records we have 10 pages in total (100/10), thus the cost to access the entire file is 10 I/O (or 10 pages)

Which alternative is better?

- **Example:** Find all records with ages between 20 and 30, for the file that has B pages. Consider both alternative: having an unsorted and sorted file. What would be the cheapest cost?
- $20 < \text{age} < 30$, num pages = B

- Heap file (no order) = B ;

11,80
12,10
52,20
13,75

14,80
14,10
36,75
18,75

20,80
21,10
22,20
16,20

36,80
41,10
12,20
80,75

Heap file

- Sorted file (exploit order) = $\log_2 B$

11,80
12,10
12,20
13,75

14,80
14,10
16,20
18,75

20,80
21,10
22,20
36,75

36,80
41,10
52,20
80,75

Sorted file

Cost of Operations (in # of I/O's)

B: Number of data pages (NOT records)

	Heap File	Sorted File	Notes
Scan all records	B <i>Average case</i>	B	
Equality Search	0.5B	$\log_2 B$	<i>Assumption (exactly one match)</i>
Range Search	B	$(\log_2 B) + (\text{\#match pages})$	
Insert	2	$(\log_2 B) + 2*(B/2)$	<i>must Read & Write</i>
Delete	$0.5B + 1$	$(\log_2 B) + 2*(B/2)$	<i>must Read & Write</i>

Heap file faster to **insert**, Sorted file faster to **search**

- Assume 10,000 pages approximately 32,000 records
- Scan all records

```
SELECT *  
FROM employee
```

Heap: 10000 I/O **Sorted:** 10000 I/O

- Equality Condition

```
SELECT *  
FROM employee  
WHERE departmentID = 5
```

Heap: 5000 **Sorted:** $\log_2 10000$ (13.2) 14 I/O #

- Range Search

```
SELECT *  
FROM employee  
WHERE salary > 85000  
AND salary < 100000
```

Heap: 10000 **Sorted:** $\log_2 10000$ (13.28) 14 + match pages (e.g 588**) = 602 I/O

We must take the ceiling value i.e.next highest integer **588 assumes approximately 5.9% of all records fall into this range)

Examples of Range Conditions

```
1  SELECT * FROM t1  
2    WHERE key_col > 1  
3    AND key_col < 10;  
4  
5  SELECT * FROM t1  
6    WHERE key_col = 1  
7    OR key_col IN (15,18,20);  
8  
9  SELECT * FROM t1  
10   WHERE key_col LIKE 'ab%'  
11   OR key_col BETWEEN 'bar' AND 'foo';
```

INFO90002

- Assume 10,000 pages approximately 32,000 records
- INSERT

```
INSERT INTO employee  
VALUES (32983, 'Andrew', 'Dale', 43000, 11, 1, '1975-11-10');
```

Heap: 2 **Sorted:** $\log_2 10000 (13.28) 14 + 2 * (10000/2) = 10014$

- DELETE

```
DELETE FROM employee  
WHERE employeeID = 15;
```

Heap: 5001 **Sorted:** $\log_2 10000 (13.28) 14 + 2 * (10000/2) = 10014$

Heap file faster to INSERT, Sorted file faster to SEARCH

DATA STORAGE

- File organization (Heap & sorted files)
- Index files & indexes
- Index classification



THE UNIVERSITY OF
MELBOURNE

Indexes

Index Types

Clustered v Unclustered
Primary v Secondary
Single Key , Composite Key
Tree based , Hash based

Indexing

- Sometimes, we want to retrieve records by specifying the *values in one or more fields*, e.g.,
 - Find all students in the “**CIS**” department
 - Find all students with a **gpa** > 3
- An **index** is a data structure built on top of data pages used for efficient search. The index is built over specific fields called **search key fields**. E.g. we can build an index on **GPA**, or department name.
 - The index speeds up selections on the *search key fields*
 - **Any** subset of the fields of a relation can be the search key for an index on the relation
 - **Note:** Search key is not the same as **key** (e.g., doesn't have to be unique)

Example: Simple Index on GPA

An index contains a collection of **data entries**, and supports efficient retrieval of **data records** matching a given **search condition**

$2 < \text{GPA} < 2.4$

Directory

Larger/equal than

Smaller than

Sorted data entries (on GPA)

1.2	1.7	1.8	1.9	2.2	2.4			2.7	2.7	2.9		3.2	3.3	3.3		3.6	3.8	3.9	4.0
-----	-----	-----	-----	-----	-----	--	--	-----	-----	-----	--	-----	-----	-----	--	-----	-----	-----	-----

(Index File)

(Data file)

Data Records
In Data Pages


Find results


DATA STORAGE

- File organization (Heap & sorted files)
- Index files & indexes
- Index classification


- Classification based on various factors:
 - Clustered vs. Unclustered
 - Primary vs. Secondary
 - Single Key vs. Composite
 - Indexing technique:
 - Tree-based, hash-based, other

- **Clustered vs. unclustered:** If order of **data records** is the same as the order of **index data entries**, then the index is called **clustered index**. Otherwise is **unclustered**.





Zoom in Unclustered Index

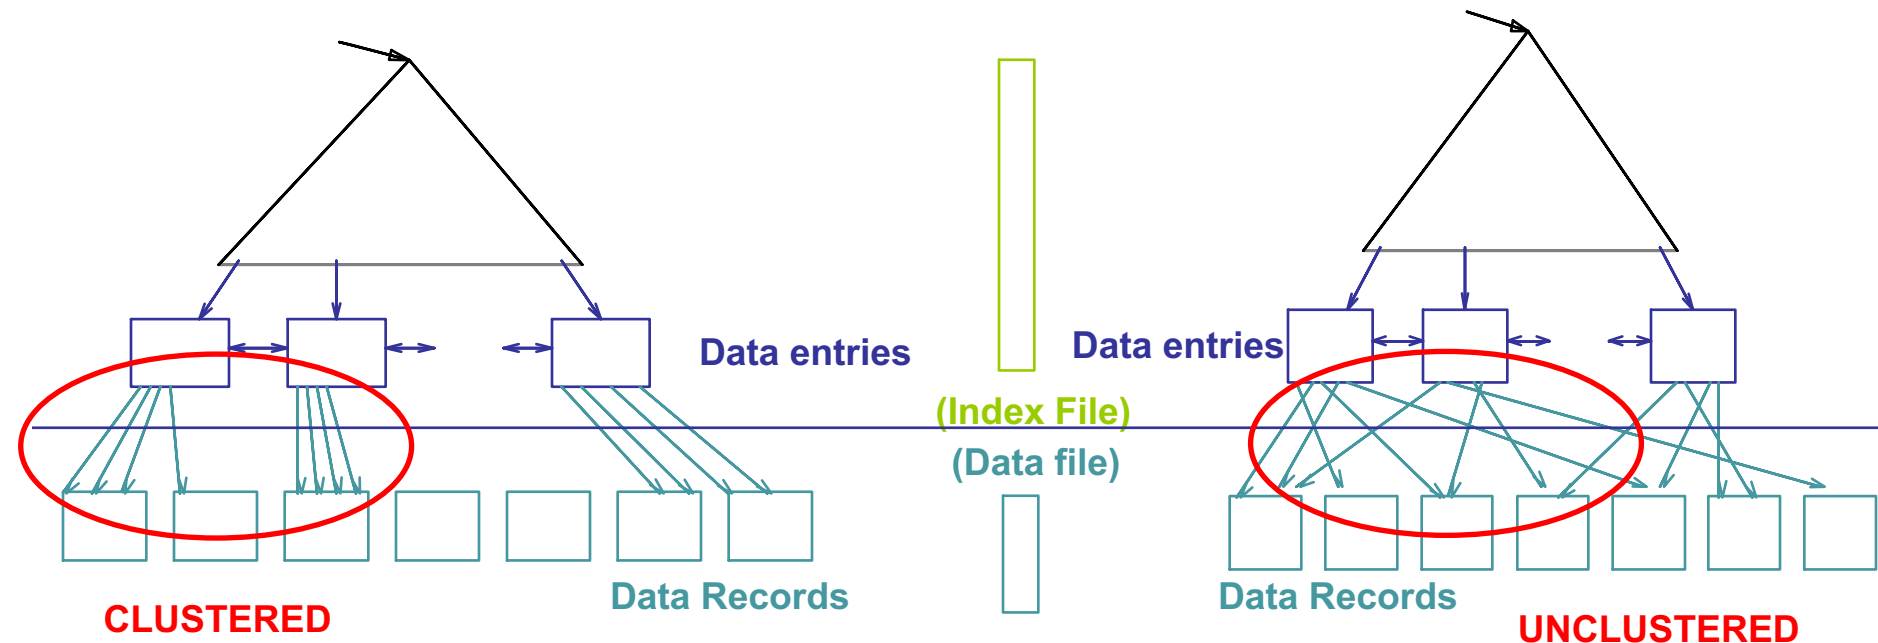


INDEXES

- A data file can have a clustered index on at most **one search key combination** (i.e. we cannot have multiple clustered indexes over a single table)
- Cost of retrieving data records through index varies *greatly* based on whether index is clustered (cheaper for clustered)
- Clustered indexes are more **expensive to maintain** (require file reorganization), but are really efficient for **range search**

Clustered vs. Unclustered Index: Cost

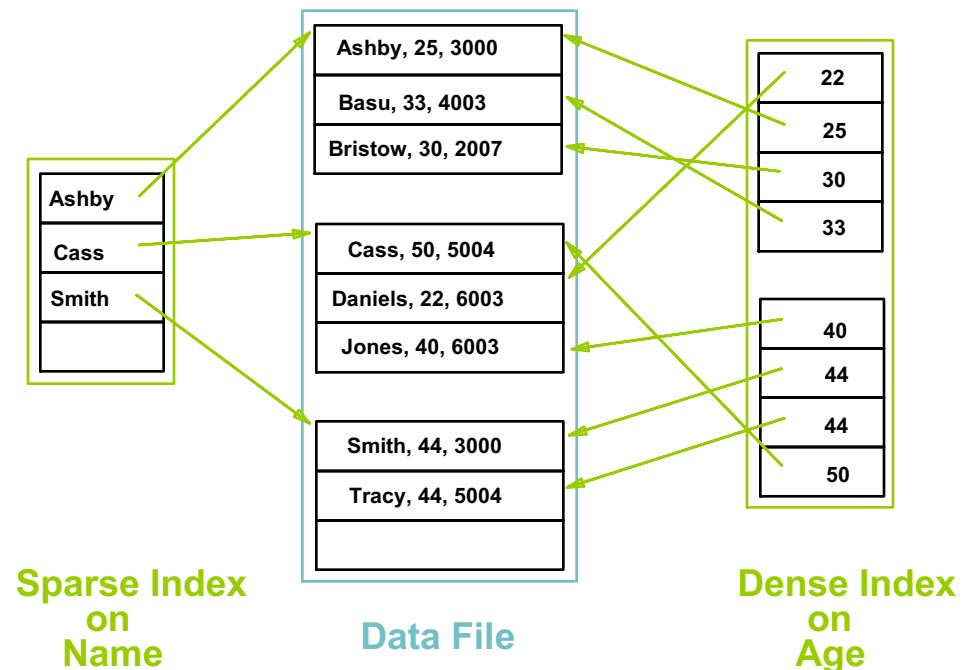
- (Approximated) cost of retrieving records found in range scan:
 1. Clustered: cost \approx # **pages** in data file with matching records
 2. Unclustered: cost \approx # of matching index **data entries (data records)**



- **Primary** index includes the table's **primary key**
- **Secondary** is any other index
- Properties:
 - Primary index **never** contains duplicates
 - Secondary index **may** contain duplicates

Dense vs. Sparse Index

- **Dense:** at least one data entry per key value
- **Sparse:** an entry per data page in file
 - Every sparse index is clustered!
 - Sparse indexes are smaller; however, some useful optimizations are based on dense indexes.
 - Alternative 1 always leads to dense index.




Composite Search Keys

- An index can be built over a combination of search keys
- Data entries in index sorted by search keys

- Examples:

1. Index on $\langle \text{age}, \text{sal} \rangle$
2. Index on $\langle \text{sal}, \text{age} \rangle$
3. Efficient to answer:
age=12 and sal = 10
age=12 and sal > 15



- Hash-based index:

- Represents index as a collection of *buckets*. Hash function maps the search key to the corresponding bucket.

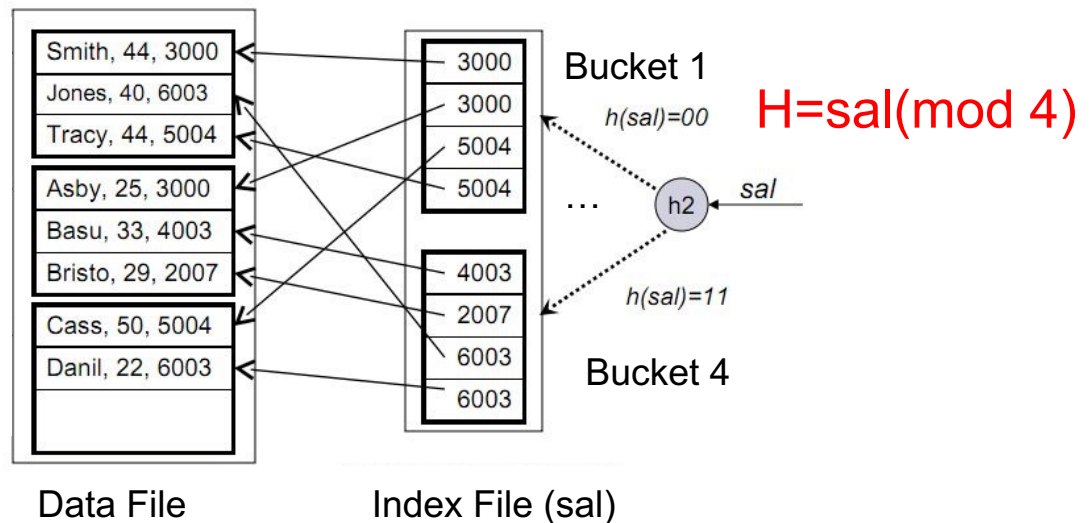
- $h(r.search_key) = \text{bucket in which record } r \text{ belongs}$

- Good for **equality** selections

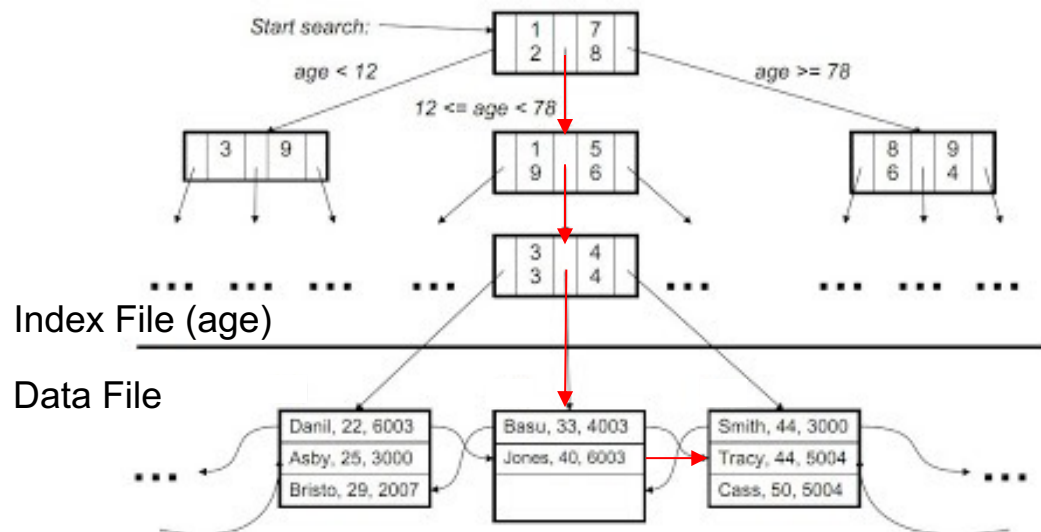
- Example:** Hash-based index on (sal)

Find Sal = 2007

$2007 \bmod 4 = 3$ go to Buck.3



- Tree-based index:
 - Underlying data structure is a binary (B+) tree. Nodes contain pointers to lower levels (search left for lower, right for higher). Leaves contain data entries sorted by search key values.
 - Good for **range** selections
 - So far we have shown those
- **Example:** Tree-based index on (age)



Find age > 39



- Many alternative file organizations exist, each appropriate in some situation
- If selection queries are frequent, sorting the file or building an *index* is important
- Index is an additional data structure (i.e. file) introduced to quickly find entries with given key values
 - Hash-based indexes only good for equality search
 - Sorted files and tree-based indexes best for range search; also good for equality search
 - Files rarely kept sorted in practice (because of the cost of maintaining them); B+ tree index is better



INFO90002

- Describe alternative file organizations
- What is an index, when do we use them
- Index classification

* All material is examinable – these are the suggested key skills you would need to demonstrate in an exam scenario



- DBMS Architecture & Distributed Databases