



# INFO90002 Database Systems & Information Modelling


Lecture 15: Database Architecture & Administration

MELBOURNE

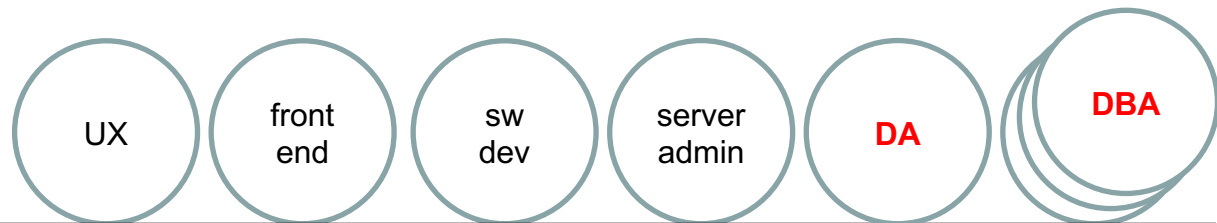
- The 'Database Administrator' role
- Architecture – Understanding the DBMS
  - concepts
- Performance improvement
  - concepts
  - common approaches e.g. indexes

# The DBA role

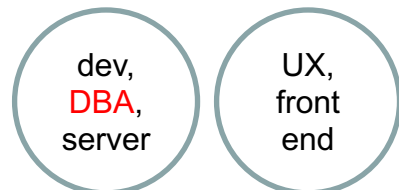
- primarily concerned with “maintenance” / “ops” phase
- but should be consulted during all phases of development
- “Database Administrator” or “DBA”
- often framed as a “job” or a “person”
- Large companies – many DBA’s
- Small company developer is the DBA
- DBA role can be made redundant by Cloud-based DBMS or “database as a service” DAAS (often IAAS or PAAS)



large org



small org



MELBOURNE

- **Data Administrator** (CDO / CSO) (management role)
  - data policies, procedures and standards
  - planning
  - data conflict resolution
  - managing info repository
  - internal marketing & education
  - Compliance with legislation (EU GDPR AUS Privacy Act)
  - Compliance with company policy (e.g. Unimelb privacy policy)
- **Database Administrator** (technical role)
  - analyze and design DB
  - select DBMS / tools / vendor
  - install and upgrade DBMS
  - tune DBMS performance
  - manage security, privacy, integrity
  - backup and recovery



THE UNIVERSITY OF  
MELBOURNE

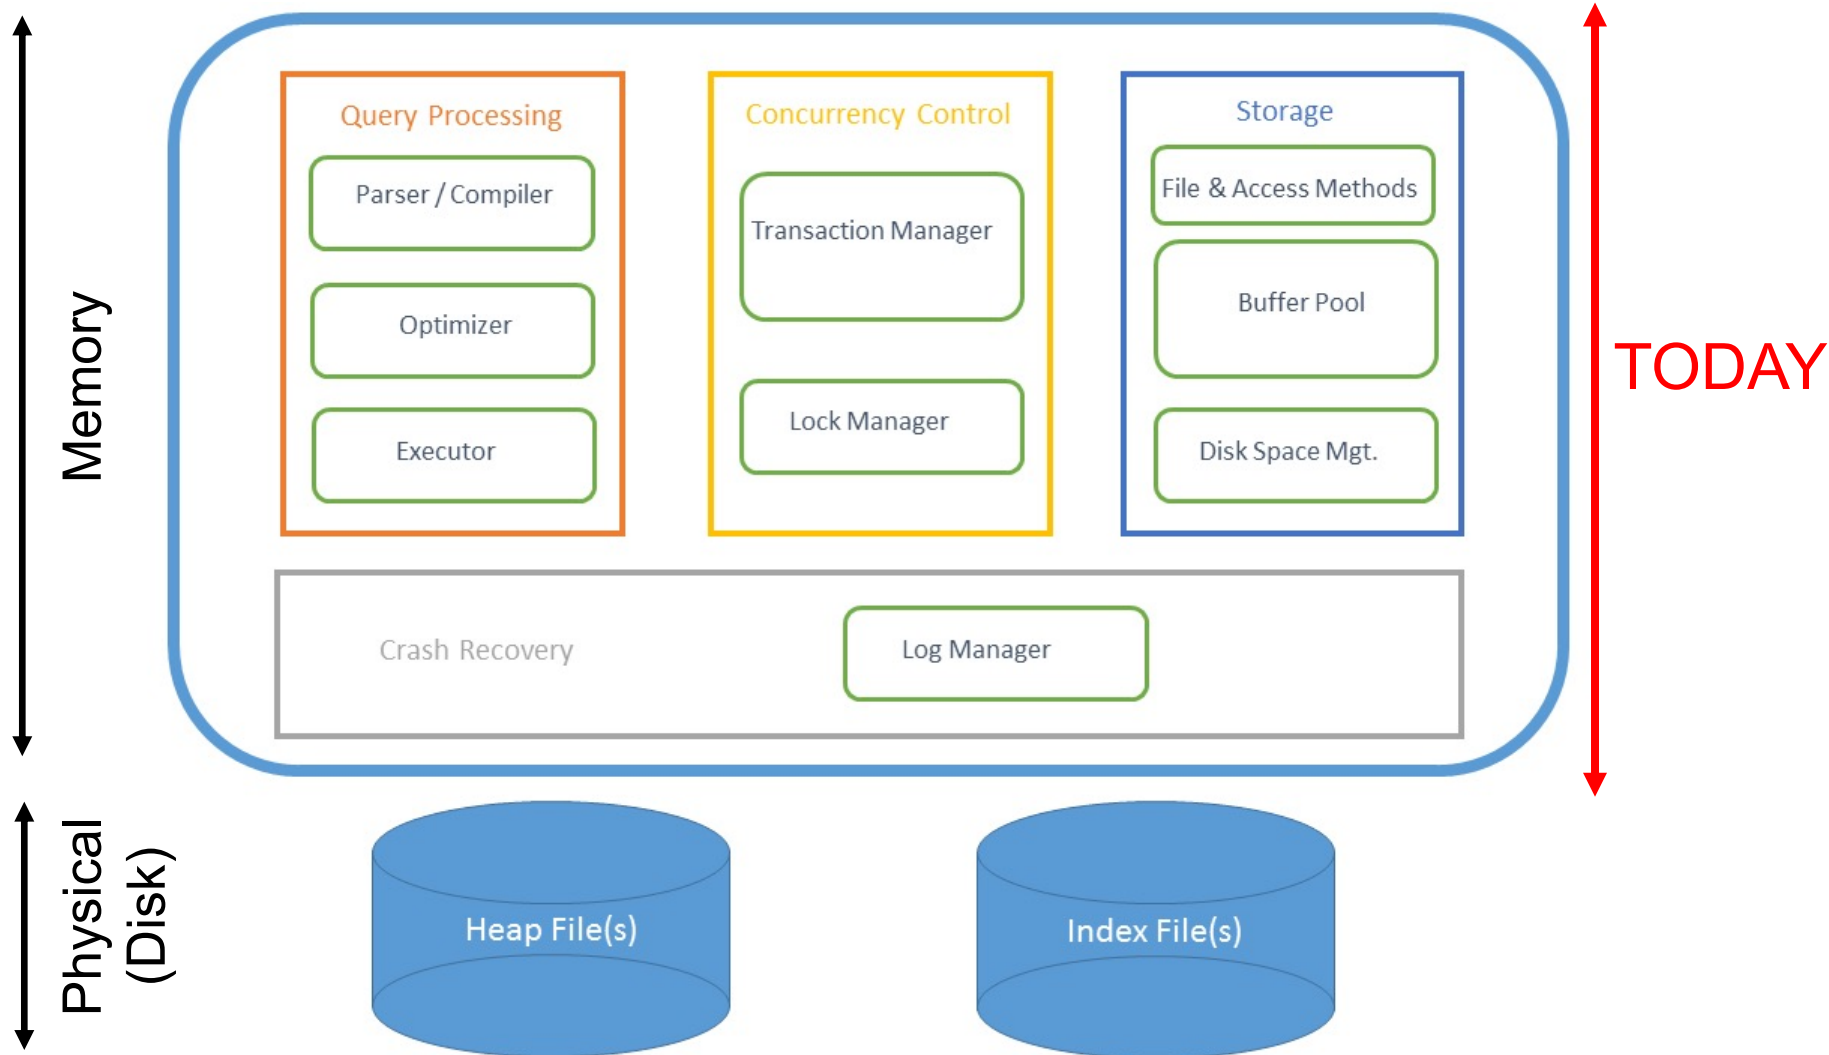
# Architecture of a Database Management System (DBMS)



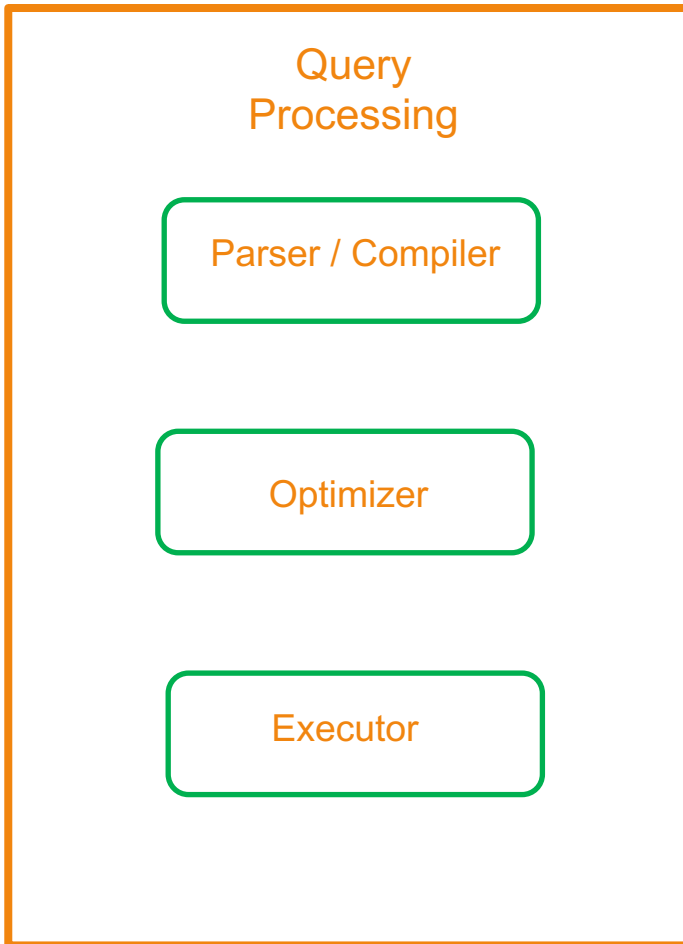
MELBOURNE

- A Database Management System (DBMS)
- Exists as one entity in two places
  - In Memory
  - Physically on disk
- Both places manage
  - Data (the reason we have the DBMS)
  - Performance (how it performs as it is used & grows)
  - Concurrency (manages high volumes of users)
  - Recoverability (assist in recovery and availability)
- One place is persistent the other transient
  - Disk representation is always present
  - Memory – transient – only exists when DBMS is running

MELBOURNE



Graphic Courtesy of  
Dr Renata Borovica-Gajic



## Parsing

- Syntax is correct & can “compile”
- DBMS User Permissions
- Resources (Data, Code, be able to Record Changes/Retrieve results)

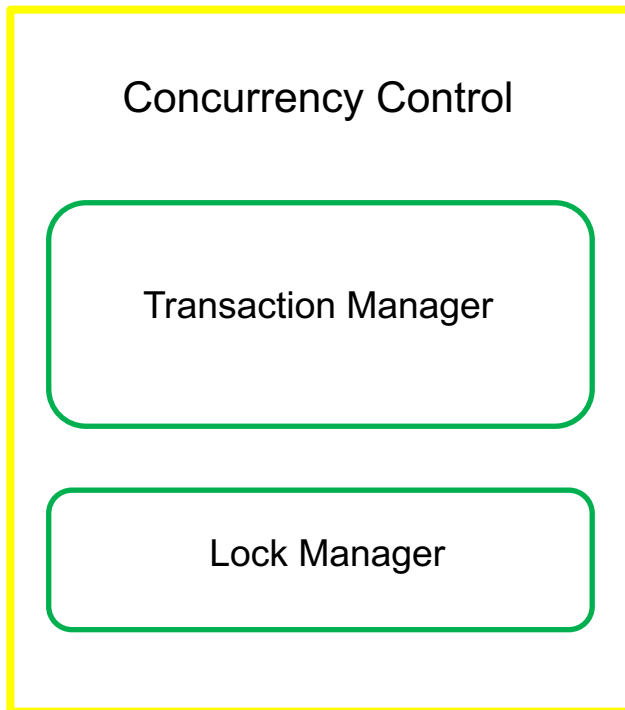
## Optimizing

- Execution Plan & Execution Cost
- Evaluate indexes, table scans, hashing
- Eliminate worst, consider best options
- Lowest cost theoretically “best”

## Execution

- Meet the ACID test,
- Atomic: All rows succeed, or all fail
- Ensure resources are available
  - Data, Log changes, Memory, Cursor to do the work for the USER






- Manages the work of the DBMS.
- Transaction Manager handles all aspects of the SQL transaction - which DBMS user wants WHAT resource
- Lock Manager is a list of what resources are locked and by which user at what level ( & who is waiting)
- not only tables, indexes
  - buffers, cursor, memory addresses of resources

MELBOURNE

- Essential to manage large scalable DBMS
- Enables 1,000,000s of concurrent users
- Like a Traffic Policemen controlling the flow of traffic
  - Who can do what (allowed to do what they need to do)
  - Who has to wait (queue)
  - Who can travel through the intersection concurrently
    - Usually *readers of data*
- What transactions have completed, in progress, compiled.
  - What resources are involved with that transaction
  - Who last used, is using and wants those resources
    - SQL, Cursor, Index, Table, Rows, File Access, Recovery Logs



- File & Access Methods
  - Disk to Memory to Disk
  - Read a buffer or a block of buffers
- Buffer Pool
  - Data in memory
    - Row data
    - Index data
  - Organised
- Disk Space Management
  - How to organise growth of data on disk efficiently by writing efficiently.




MELBOURNE

- How to access the file
  - Full Scan (full table scan)
    - From beginning to end of the entire file
  - Partial Scan (index range)
    - Using an index to scan a range of values
      - ( 2 > deptno > 9)
  - Page only (file header and page)
  - Read the index ***and*** data file

MELBOURNE

- Buffer Pool
  - All data (including indexes) is stored in the buffer
  - May contain multiple copies of the same data
    - Need to know which copy is the current committed version
  - Organised using a double linked list (see Storage & Indexing lecture)
- Disk Space Management
  - How to allow files to grow on disk (and set max growth size)
  - File organization
    - e.g. index reorganization;
    - varchar growth e.g. Brown (5) grows to Nicholson (9) for last name

- Hierarchical Structure
- Storage of Object (tables, indexes, rollback, logs)
- Row by Row space management inefficient
- Page better
- Block Best



- Many Object Types (Tables, Indexes, Undo)
- Each buffer contains rows, b+ tree leaf etc.
- Each buffer can have one of four status types:

- Current

- In use current committed version of data (row)

- Active

- Most recent change (may not be committed)

COMMIT (Current: DepartmentID=9)

- Stale

- An old version of the data

- Aged

- Old and about to be removed from buffer pool

Buffer Pool

DepartmentID=8

DepartmentID=9

*DepartmentID=8*

*DepartmentID=8*



MELBOURNE

Crash Recovery

Log Manager

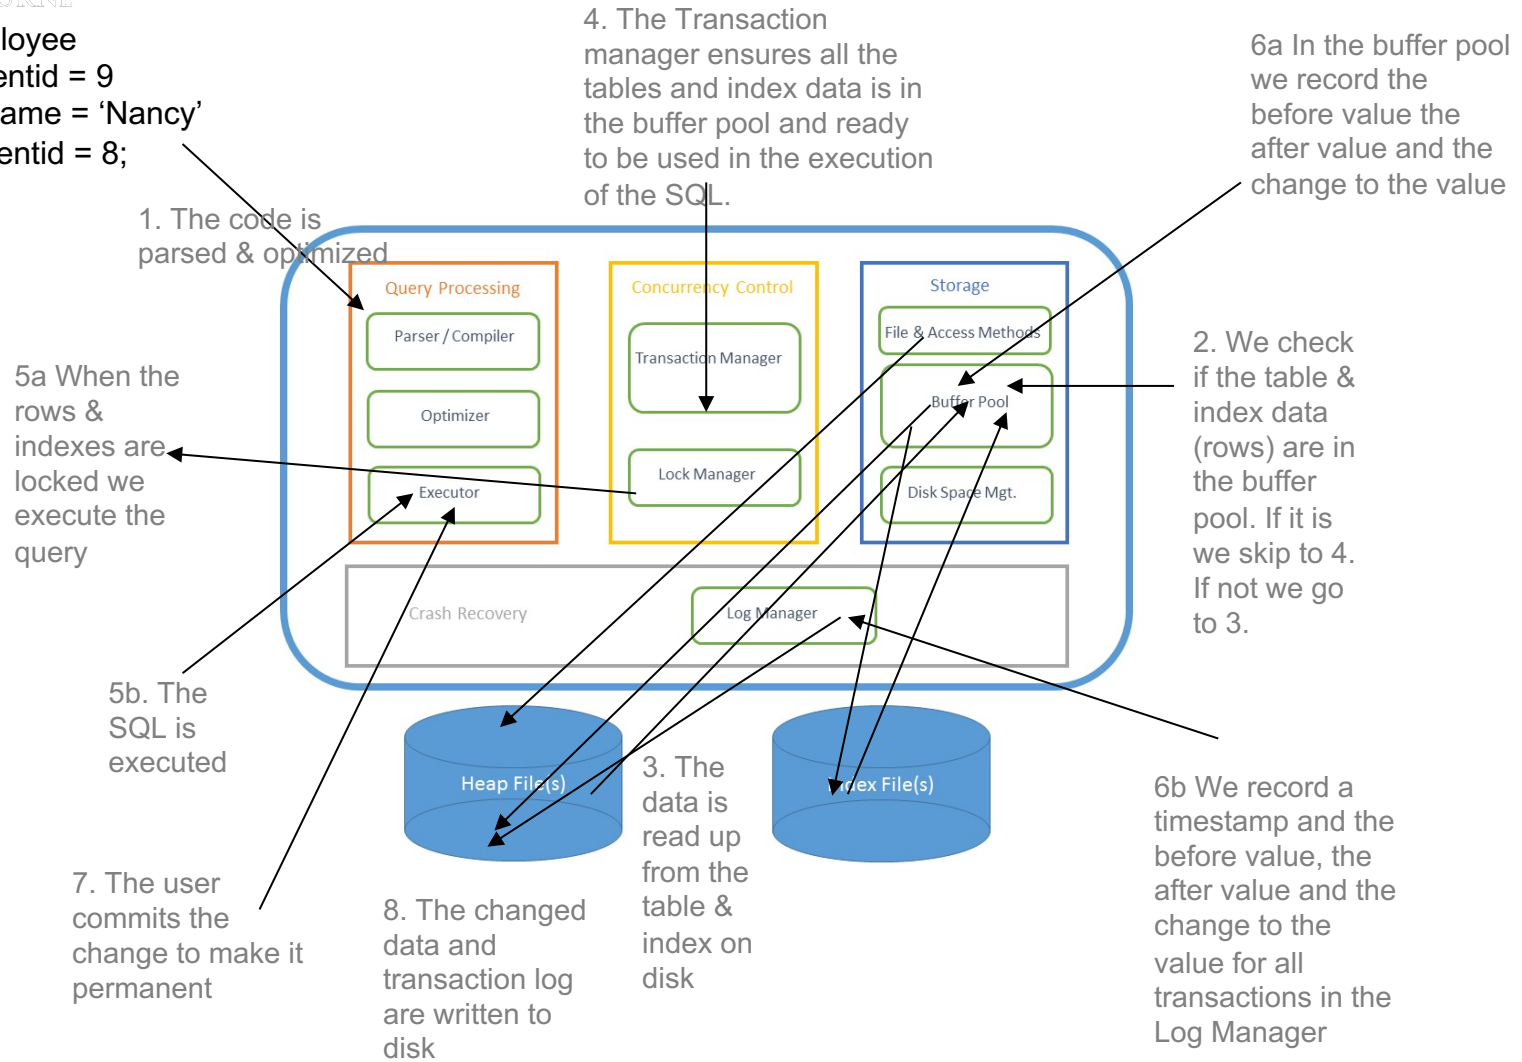
- Recovery
- Log Manager records ALL changes
  - Statement
  - Transaction
    - Statement
    - Rollback values
    - Before and After values
    - Timestamp begin
      - transaction, savepoint & commit timestamps
  - Database
    - Data Dictionary Changes



# DBMS – How a transaction works

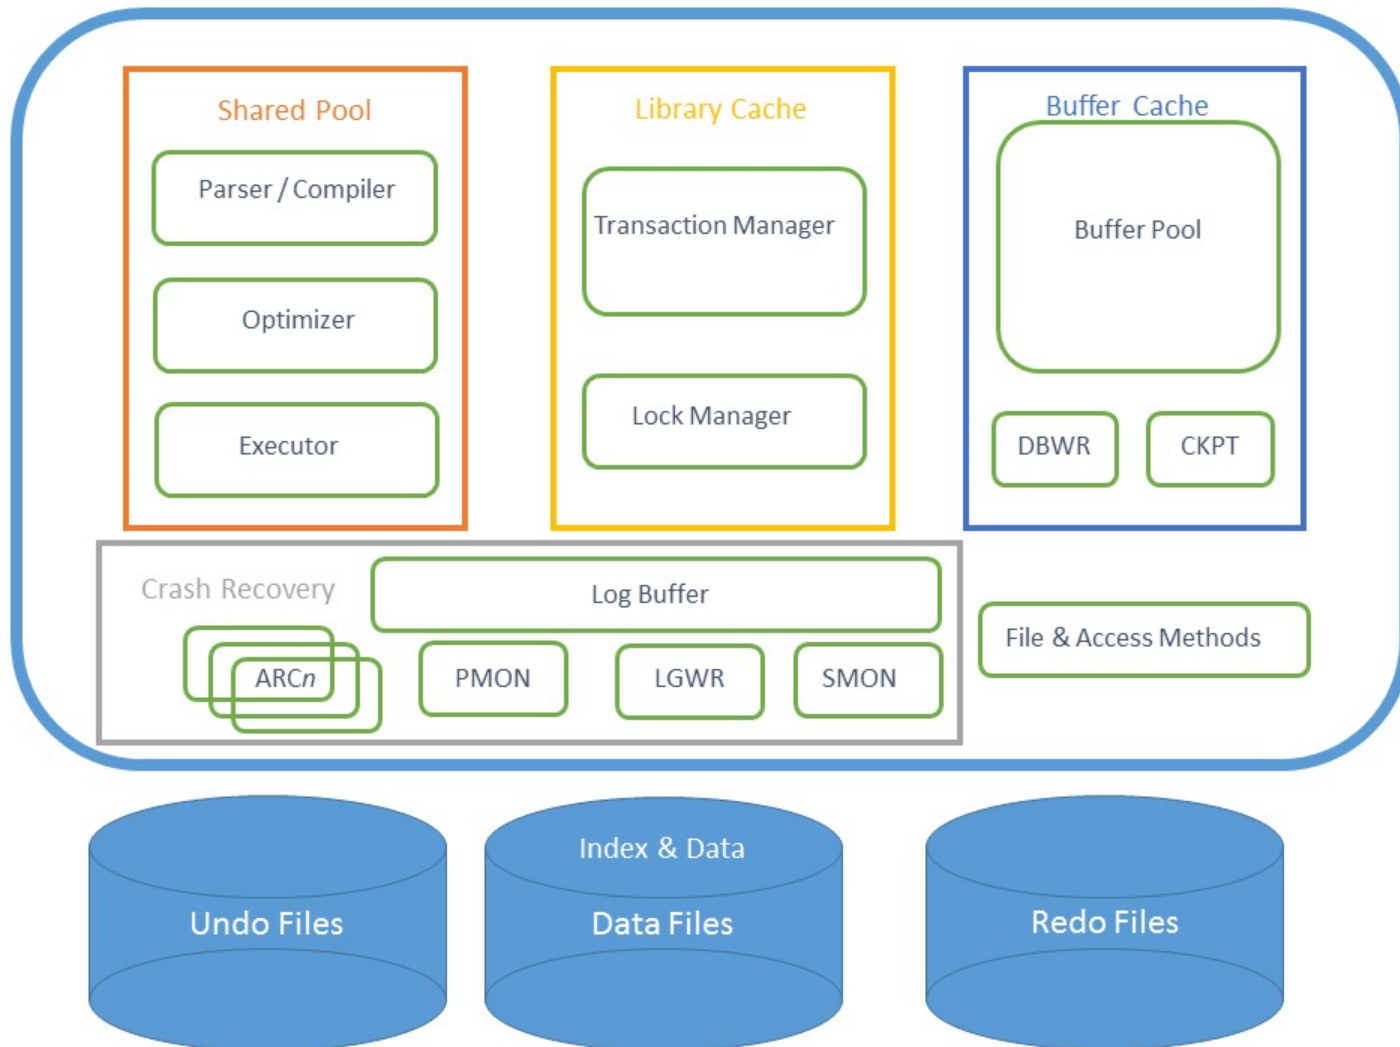
MELBOURNE

UPDATE employee  
SET departmentid = 9  
WHERE firstname = 'Nancy'  
AND departmentid = 8;



Graphic Courtesy of  
Dr Renata Borovica-Gajic

# How Oracle\* DBMS looks



**\* This slide is not examinable**

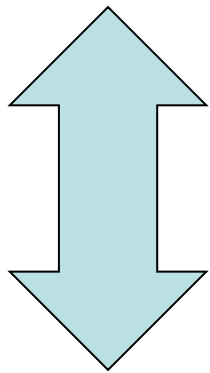


THE UNIVERSITY OF  
MELBOURNE

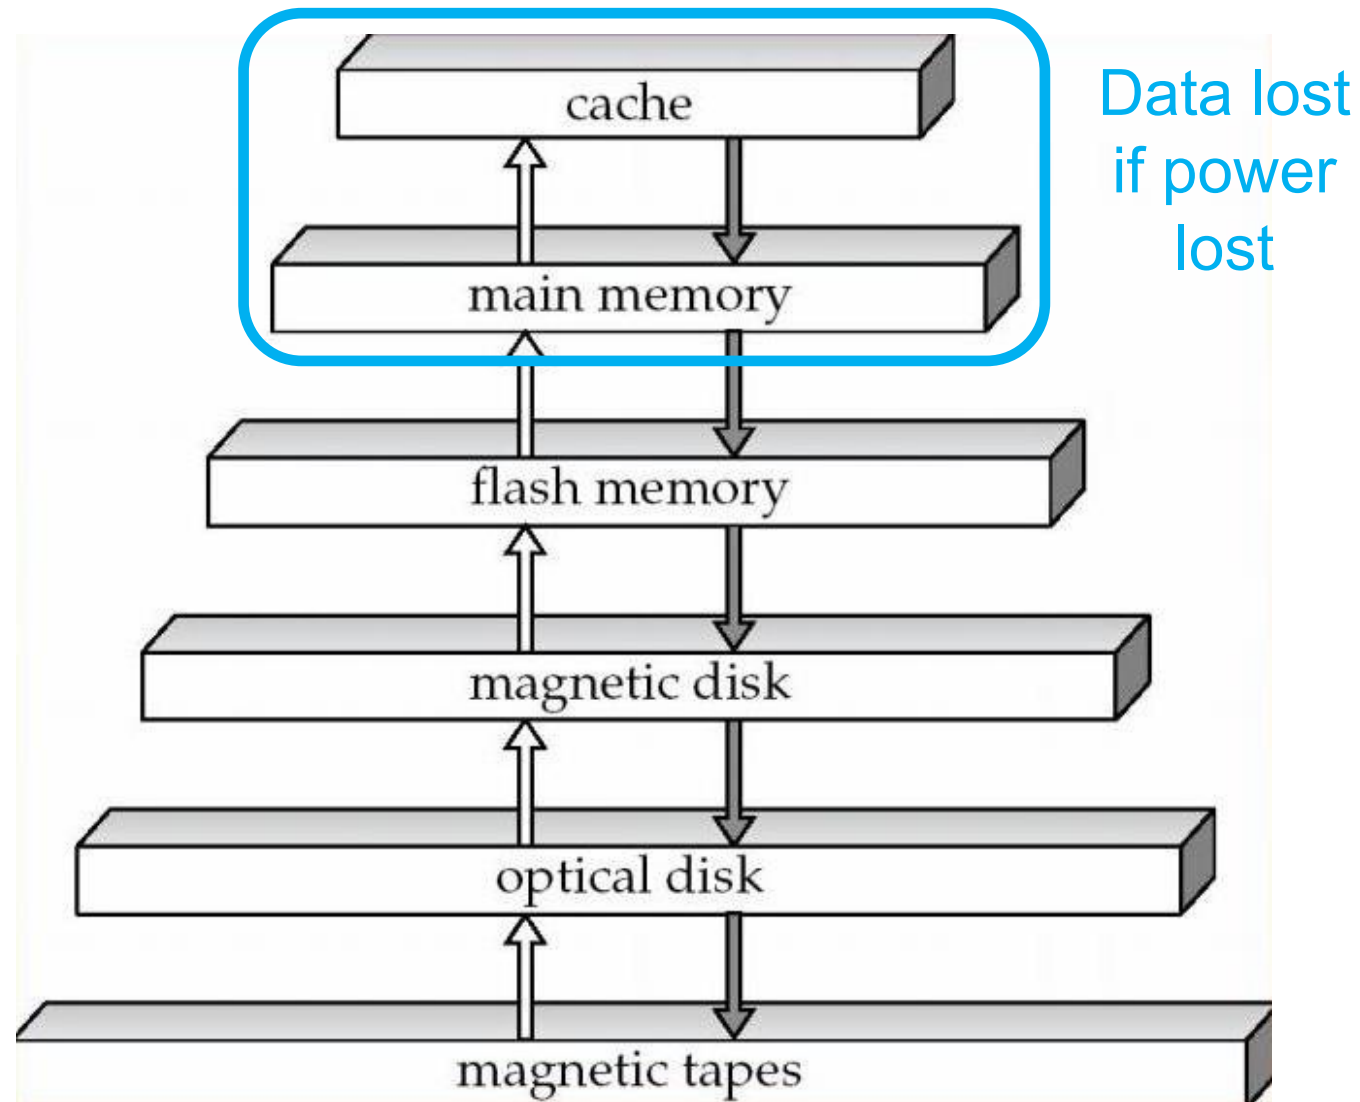
# Database Performance

# Storage media hierarchy

faster, more  
expensive,  
smaller  
capacity



slower,  
cheaper,  
*older*,  
bigger  
capacity



- Caching data in memory, e.g. data buffers
- Placement of data files across disc drives
- Fast storage such as SSD
- Database replication and server clustering
- Use of indexes to speed up searches and joins
- Good choice of data types (especially PKs)
- Good program logic (no long running CRUD)
- Good query execution plans
- Good code (no deadlocks)

- Data and Code found in memory
- Avoids a read
- Reads are expensive
- Goal in to minimize reads (& writes)
  - Writes are necessary (recovery logs, changed data)
- “in memory databases”
  - all code all data loaded into memory on db start & stays until shutdown

Buffer Pool  
(Table & Index rows)

Parser / Compiler  
(SQL)

MELBOURNE

- Spread the files across the physical server
  - RAID (0, 0 +1, 5)
- We can't avoid writes
  - Spread files across many disks
    - Avoid contention
      - (many users competing for same resource)
  - Recovery Logs (always writing)
    - faster disk
- SSD (Solid State Drives)
  - No moving parts – nothing to break down
  - Faster I/O (Input & Output compared to other disk types)

- Distributed data
  - Spreads the load
  - Data kept only where it is needed
  - Less work per physical server – faster response times



- Replicated Data
  - Spreads Load
  - Less work per physical server – faster response times





# When to create indexes

- for each table, choose the columns you will index:
- *queried frequently* (used in WHERE clauses)
- used for *joins* (PK to FK)
- primary *keys* (automatic in most DBMS)
- foreign *keys* (automatic in MySQL)
- unique columns (automatic in most DBMS)
- large tables only - small tables do not require indexes
  - if you frequently retrieve less than about 15% of the rows
- wide range of values (good for regular indexes).
- small range of values (good for bitmap/hash indexes).

covered already in week 6

source: Oracle® Database Application Developer's Guide

- Good Data Types
  - INTEGERS for PK FK & PFK (for performance)
- Good Program Logic & Code
  - Transaction design

BEGIN TRANSACTION

SELECT

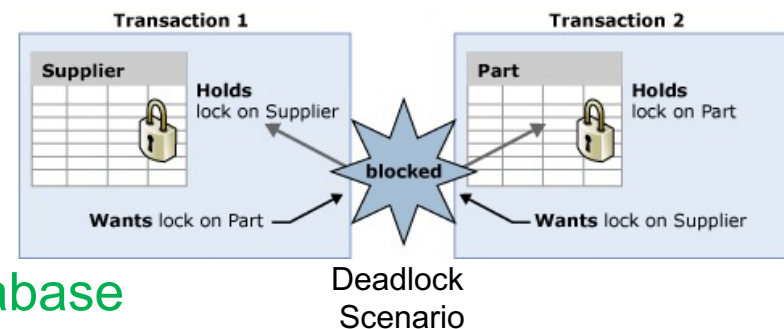
UPDATE

UPDATE

UPDATE

COMMIT

- Avoid Long Complex Transactions that never commit or savepoint
- Avoid coding deadlocks!
- Appropriate Locking strategy
  - Row b4 Buffer b4 Table b4 Database
  - Consider Lock Timeouts (if not automatic)





- The best execution plan has the lowest “cost”
- Known as Cost Based Optimization (CBO)



Index on where condition (cust\_city) improves cost:

```
SELECT cust_first_name, cust_last_name, cust_marital_status,
       cust_city, cust_income_level
FROM Customers
WHERE cust_last_name = 'Parkburg'
AND cust_first_name = 'Peter'
AND cust_city = 'Trafford';
```

	cust_first_name	cust_last_name	cust_marital_status	cust_city	cust_income_level
	Peter	Parkburg	single	Trafford	H: 150.000 - 169.999

```
CREATE INDEX cust_city_idx
ON customers(cust_city);

SELECT cust_first_name, cust_last_name, cust_marital_status,
       cust_city, cust_income_level
FROM Customers
WHERE cust_last_name = 'Parkburg'
AND cust_first_name = 'Peter'
AND cust_city = 'Trafford';
```





MELBOURNE

- What a DBA and Data Administrator do
  - And the difference in each role
- Database Architecture
  - Label all memory structures & know their role
- What affects database performance
  - Caching; Datafile placement; Fast Storage; Indexes; Data types; Query Execution plans; Efficient code;
- When to create an index

\* All material is examinable – these are the suggested key skills you would need to demonstrate in an exam scenario

MELBOURNE

- Distributed Databases



# INFO90002

## Database Systems & Information Modelling

### Lecture 15

### Database Architecture & Administration